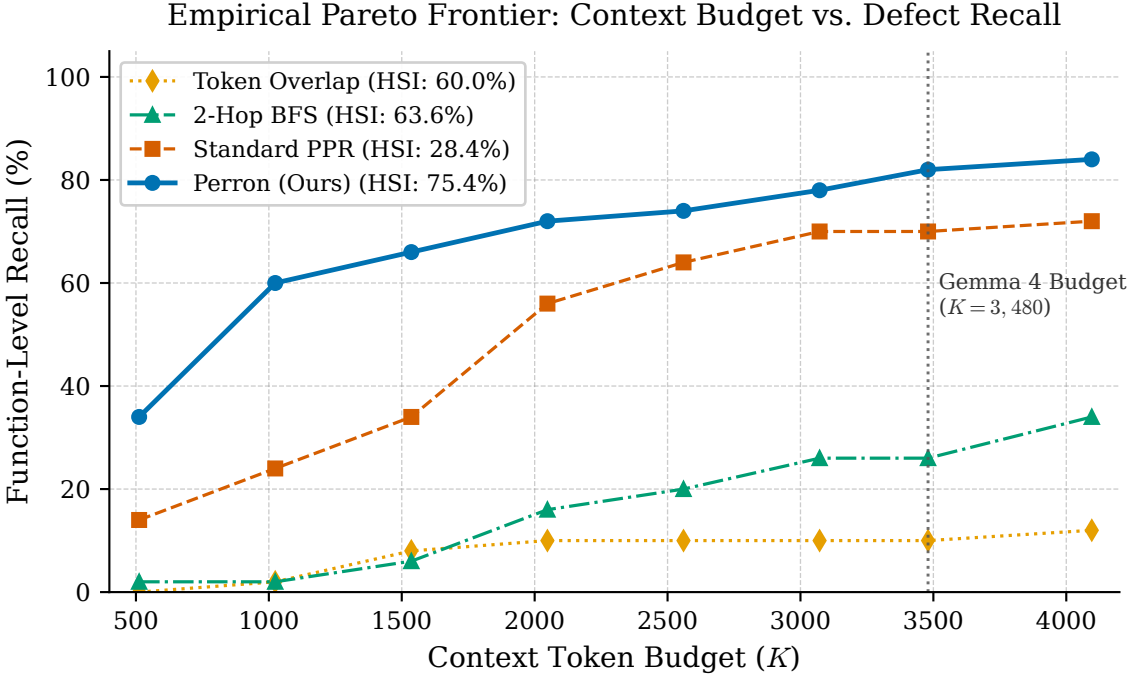
Empirical Pareto Frontier: Context Budget vs. Defect Recall
Token Overlap (HSI: 60.0%)
2-Hop BFS (HSI: 63.6%)
Standard PPR (HSI: 28.4%)
Perron (Ours) (HSI: 75.4%)
Gemma 4 Budget
($K = 3,480$)
Function-Level Recall (%)
Context Token Budget (K)
0
20
40
60
80
100
500
1000
1500
2000
2500
3000
3500
4000