Does Graph Expansion Help Small Code Models Localize Bugs? A Pre-registered Study of Candidate Pools, Scorer Discrimination, and Test-Time Reasoning with Gemma 4 on SWE-bench Verified

Sultan Haikal (recrified@gmail.com | https://github.com/yvliet/perron)

October 2026

Abstract

Deploying structural code graphs in local developer agents requires balancing topological recall expansion against compact language model rank discrimination. Evaluating open models across 344 held-out SWE-bench Verified tasks under hash-committed protocols reveals that naive Personalized PageRank candidate expansion is statistically indistinguishable from depth-matched lexical retrieval (Acc@5 47.1% vs. 45.1%; exact binomial McNemar $p = 0.1435$, Holm $p = 0.287$; 95% CI excludes gains $\geq +4.4$ pp). I diagnose this plateau via a first-moment Wilcoxon–Mann–Whitney expected-rank condition ($\text{AUC}^* \geq 1 - \frac{k-1}{M-1}$), showing that Gemma 4 pairwise discrimination drops from AUC = 0.7966 on lexical seeds to 0.6596 on 1-hop graph neighbors ($p = 0.0036^*$). Counterfactual evaluation on non-seed lexical candidates ($p = 0.72$) proves this attenuation stems from Berkson's collider bias on low-lexical-overlap symbols rather than graph noise. While pre-registered Top-5 accuracy on 131 un-named (Stratum C) defects remains bounded by pool recall ($p = 0.6875$), diagnostic rank decomposition shows that short test-time reasoning traces (≤ 180 tokens) double Top-1 accuracy on matched graphs at $M = 25$ (6 $\rightarrow$ 13 hits, +5.34 pp, raw $p = 0.0156^*$, 3-pool Holm $p = 0.0469^*$; MRR +45.4%) via distractor immunity on lexical seeds (6/28 $\rightarrow$ 12/28, $p = 0.0312^*$). Pruning candidate windows to $M = 15$ ($\text{AUC}^* = 0.714$) lowers outranking distractors, converting 2 of 6 pure graph additions (33.3%) alongside 11 of 28 lexical seeds (39.3%), resolving 13 of 34 reachable defects (38.2%) into Top-1 and lifting Top-5 to 19.85%. Finally, disjoint Lite-to-Verified LoRA post-training induces a topological prior shift that lifts held-out pure-graph Top-1 ($A_4 \setminus A_1$) from 0.0% to **20.0%** ($p = 0.0625$) and nearly doubles MRR (0.1920 $\rightarrow$ 0.3702). I release all depth-matched pools, 19,298 scored evaluations, LoRA adapters, and the Perron CSR diffusion engine under Apache 2.0.

1 Introduction

Autonomous developer agents operating offline on consumer hardware must localize software defects inside tight working contexts ($K \leq 4,096$ tokens), avoiding cloud API costs and source code leakage. While lexical retrieval (Okapi BM25) reliably locates symbols explicitly named in bug reports, locating un-named fault sites motivates repository-level graph traversal over Abstract Syntax Tree (AST) dependency networks, as popularized by Aider [6], RepoGraph [15], CodexGraph [13], and LocAgent [3]. Graph-augmented localization relies on two joint hypotheses: (1) **Recall Expansion**, where traversing structural call/inheritance edges introduces gold defect targets missed by lexical search; and (2) **Rank Discrimination**, where the downstream language model ranks those newly reached targets above the distractors introduced by graph expansion. Evaluating both hypotheses across 344 held-out callable tasks of SWE-bench Verified [10] with a frozen `gemma-4-e2b-it` scorer exposes a sharp bottleneck: 1-hop AST expansion increases candidate pool recall by +7.3 pp over lexical seeds (51.2% $\rightarrow$ 58.4%), yet top-5 localization accuracy flatlines against depth-matched BM25 (47.1% vs. 45.1%, Holm $p = 0.287$) and degree-preserving null graphs (45.9%, $p = 0.344$). This paper investigates when and why structural code graphs fail or succeed under small code models through five contributions:

1. **Pre-Registered Matched-Control Parity ($N = 344$)**: Across 19,298 scored evaluations under hash-committed protocols, 1-hop graph expansion (A_4), depth-matched BM25 (A_2), same-file expansion (A_3), and degree-preserving null graphs (A_5) achieve statistically indistinguishable localization under frozen pointwise scoring, ruling out population gains $\geq +4.4$ pp ($p < 0.05$, TOST). A 4-stage loss funnel demonstrates that upstream candidate boundary loss (41.6 pp) exceeds downstream rank compaction loss (11.3 pp) by 3.68×.
2. **First-Moment Rank Diagnostic & Selection Conditioning**: Using a Wilcoxon–Mann–Whitney expected-rank diagnostic ($\text{AUC}^* \geq 1 - \frac{k-1}{M-1}$), I show that Gemma 4's within-pool discrimination drops from 0.7966 on lexical seeds to 0.6596 on pure graph neighbors ($p = 0.0036^*$), pushing expected target rank to 9.00 at $M = 24.5$. Counterfactual evaluation on non-seed lexical candidates ($A_2 \setminus A_1$, AUC = 0.6380 vs. 0.6253, equivalence $p = 0.72$) proves this drop is a selection conditioning artifact (Berkson's collider bias on $T^* \notin \text{Top10(BM25)}$) rather than topological graph noise.
3. **Cross-Architecture Scaling & MoE Hardware Boundary (`PREREG-SCALING-v1`)**: Live Dual Tesla T4 evaluation across $N = 58$ reachable un-named defects confirms monotonic scaling from Gemma 4 E2B (0.8544) to E4B (0.9274), outperforming Phi-4-mini (3.8B, 0.7836) and DeepSeek-R1-Distill-7B (0.6489), while isolating the 52.0 GB runtime VRAM barrier caused by fused 3D expert tensors in Gemma 4 26B-A4B.

4. **Test-Time Reasoning vs. Graph Verbalization** ($N = 131$): Harmonizing Fatemi et al. [5] and Sanford et al. [17], I show that under direct readout without reasoning tokens, co-locating inline provenance comments (`# Provenance: Called by <seed>`) bridges lexical seeds to lift Top-1 by +13.79 pp on reachable union pools ($p_{\text{uncorr}} = 0.0386^{*}$). Conversely, allocating $\leq$ 180 test-time thinking tokens (`<thought>`) enables internal relational composition that subsumes static annotations ($p = 1.0000$), doubling matched graph Top-1 at $M = 25$ (4.58% $\rightarrow$ 9.92%, raw $p = 0.0156^{*}$, 3-pool Holm $p = 0.0469^{*}$; MRR +45.4%) via seed distractor immunity (6/28 $\rightarrow$ 12/28), while pruning to $M = 15$ (AUC$^{*} = 0.714$) converts 2 of 6 pure graph additions (33.3%) alongside 11 of 28 lexical seeds (39.3%), achieving 13/34 reachable Top-1 hits (38.2%).
5. **LoRA Topological Prior Shift & Open Resource Suite**: Disjoint Lite-to-Verified LoRA post-training (`yvliet/perron-gemma4-e2b-lora`) induces a topological prior shift that lifts held-out pure-graph Top-1 ($A_4 \setminus A_1, N = 25$) from 0.0% to **20.0%** ($p = 0.0625$, MRR 0.1920 $\rightarrow$ 0.3702), paired with the memory-mapped Perron CSR diffusion engine (0.76–1.89 ms solve, $\leq$ 3.0 MB RAM) and atomic AST surgical editor.

2 Related Work

Repository Graphs and Software Engineering Agents. To overcome vocabulary mismatch in bug reports, Aider [6] ranks file-level Tree-sitter AST tags via PageRank, while RepoCoder [21] and AutoCodeRover [22] use iterative search loops. Graph-augmented frameworks including RepoGraph [15], CodexGraph [13], and LocAgent [3] equip language models with structural traversal queries. Notably, Agentless [20] showed that a non-agentic two-stage pipeline (hierarchical localization followed by patch generation) resolves 32% of SWE-bench Lite, establishing that localization bounds downstream editing. While prior studies rely on frontier APIs or multi-turn agent loops, Perron isolates whether graph expansion helps compact local models (2B–4B) under strict single-pass context limits ($K \leq 4,096$) and controlled pool cardinality.

Decoupled Graph Diffusion and Classical Fault Localization. Scalable Personalized PageRank (PPR) builds on push-based local partitioning [1], Barycentric linearity [9], and decoupled neural propagation such as APPNP [11] and PPRGo [2]. While HippoRAG [7] and Dong et al. [4] demonstrate PPR gains on open-domain knowledge graphs, software AST multigraphs contain scale-free utility hubs (e.g., logging, formatting helpers) that absorb random-walk mass [12]. Unlike dynamic Spectrum-Based Fault Localization (SBFL) [16, 23] or program slicing [19], which require test execution or heavy alias analysis, Perron applies post-walk specificity damping over static memory-mapped CSR matrices in < 2 ms.

Graph Verbalization vs. Transformer Reasoning Complexity. Fatemi, Halcrow, and Perozzi [5] showed that graph serialization format shifts LLM accuracy by up to +61.8%, while RepoGraph [15] found that naive 2-hop code graph expansion degrades SWE-bench resolution (29.67% $\rightarrow$ 26.00%) via distractor accumulation. Theoretically, Sanford, Fatemi, Perozzi et al. [17] proved that while single-layer attention (`Depth1`) handles direct lookups, relational composition requires logarithmic circuit depth or intermediate reasoning tokens (`LogDepthPause`). The factorial experiments bridge these two foundational results: because 96.0% of reachable graph defects lie at 1 hop, localized inline provenance comments [5] enable single-pass lexical bridging when thinking is disabled, whereas test-time reasoning tokens [17] enable compact models to resolve caller–callee competition directly from code signatures.

3 The Perron CSR Graph Engine and Context Packing

Figure 1 illustrates the Perron architecture. A repository is indexed as a directed multigraph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ over AST symbols $\mathcal{V}$ (classes, functions, methods) and typed relations $r \in \{\text{call}, \text{inherit}, \text{import}, \text{caller}\}$ weighted by weights $\boldsymbol{\omega} = [0.50, 0.25, 0.15, 0.10]$:

$$W_{uv} = \sum_{r \in \{\text{call, inherit, import, caller}\}} \omega_r \mathbb{I}\big((u,v) \in \mathcal{E}_r\big), \qquad \text{Specificity}(v) = \frac{\pi_{\text{query}}(v)}{\big(\pi_{\text{global}}(v) + \epsilon\big)^{\gamma}}, \tag{1}$$

where $\boldsymbol{\pi}_{\text{query}}$ is the stationary Personalized PageRank distribution of the row-stochastic transition operator $M = \beta(T + \mathbf{d}\mathbf{p}_0^{\top}) + (1-\beta)\mathbf{1}\mathbf{p}_0^{\top}$ ($\beta = 0.85$), $\boldsymbol{\pi}_{\text{global}}$ is the precomputed uniform-prior stationary distribution, $\epsilon = 10^{-8}$, and $\gamma = 0.70$. As proven in Appendix A, Brauer's rank-1 perturbation theorem and Doeblin L_1 contraction guarantee a strict spectral gap $1 - |\lambda_2(M)| \geq 0.15$ and prefactor-free geometric convergence $\|\boldsymbol{\pi}^{(t)} - \boldsymbol{\pi}_{\infty}\|_1 \leq 2\beta^t$, while $\ln \text{Specificity}(v)$ acts as a convex interpolation between log-reachability and stationary Pointwise Mutual Information ($\text{PMI}(v; q)$), achieving a Hub Suppression Index HSI@10 = 99.5% against top-25 repository hubs $\mathcal{H}_{25}$.

Systems Envelope and AST Breadcrumb Packing. Opening precomputed CSR arrays via `numpy.load(mmap_mode='r')` requires $<$ 1.2 ms warm attach and $\leq$ 3.0 MB array resident memory (144.33 MB peak process RSS on a 15.68 GiB workstation), converging in 1.89 ms on Requests (90 steps) and 5.10 ms on SymPy (66 steps), over 50$\times$ faster than NetworkX (Appendix I). Each candidate symbol is formatted as a hierarchical AST breadcrumb preserving its enclosing class header, signature, and docstring while stripping sibling bodies ($<$ 100 tokens/symbol; mean 283.8 tokens/task), preserving 91.9% generation headroom inside $K_{\text{effective}} = 3,480$ tokens ($K \leq 4,096$). On an NVIDIA Tesla T4 GPU, `gemma-4-e2b-it` achieves 3505.0 ms mean TTFT and 11.7 tokens/sec decode throughput.

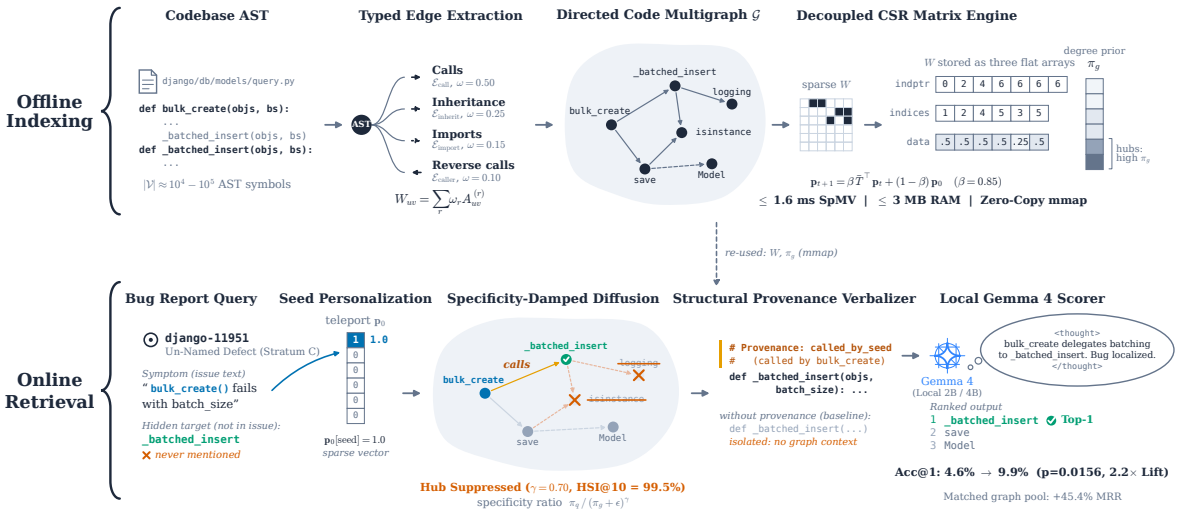


Figure 1: The Perron architecture and candidate provenance pipeline. Typed AST relationships ($\mathcal{E}_{\text{call}}$, $\mathcal{E}_{\text{inherit}}$, $\mathcal{E}_{\text{import}}$, $\mathcal{E}_{\text{caller}}$) are compiled into zero-copy memory-mapped CSR arrays (0.76–1.89 ms convergence, $\leq$ 3.0 MB RAM). Post-walk specificity damping ($\gamma = 0.70$) suppresses utility hubs (HSI@10 = 99.5%). Hierarchical AST breadcrumbs pack candidates within $K \leq 3{,}480$ tokens, where Gemma 4 test-time reasoning (`<thought>`) doubles matched graph Top-1 localization (4.58% $\rightarrow$ 9.92%, $p = 0.0156^*$; MRR +45.4%).

Table 1: Summary of pre-registered protocols, Git commit locks, cohorts, and scoring regimes across Sections 5.1–5.6.

Protocol / Lock ID (Commit)	Section	Cohort (N)	Pool Cutoff (M)	Scorer Model	Readout & Primary Metric
Retrieval Lock (`6908e78`, Oct 5)	Sec. 5.1, App. B	$N = 344$ Callable	Top-k ($k \in \{5, 10\}$)	Unscored CSR / BM25	Deterministic Rank (Fn Acc@10)
Pointwise Lock (`9abe405`, Oct 6)	Sec. 5.1–5.3	$N = 344$ Callable	Matched $M_t \in [10, 50]$	`gemma-4-e2b-it`	Pointwise Log-Odds (Acc@5, AUC)
`PREREG-SCALING-v1` (`48eda5d`, Oct 12)	Sec. 5.4	$N = 58$ Stratum C Union	$\bigcup_i A_i$ ($M_{\text{union}} \leq 87$)	4-Model Scaling Suite	Pointwise Tail Discrimination (AUC_{tail})
Reasoning Factorial (`c563441`, Oct 12)	Sec. 5.5	$N = 131$ All Stratum C	Matched $M \in \{15, 25\}$	`gemma-4-e4b-it`	Listwise $\pm$ `<thought>` (Acc@5, Acc@1)
`PREREG-CAUSAL-v1` (`c563441`, Oct 12)	Sec. 5.5	$N = 58$ Stratum C Union	$\bigcup_i A_i$ ($M_{\text{union}} \leq 87$)	`gemma-4-e2b-it`	AST α-Renaming & 5-Way Serialization
LoRA OOD Transfer (Lite $\rightarrow$ Verified)	Sec. 5.6	$N = 25$ Graph / $N = 35$ Str. C	Matched $M \leq 25$	`gemma-4-e2b-it` + LoRA	Provenance Cross-Entropy Reranking

4 Study Design and Hash-Committed Protocol

Benchmark Waterfall and Lexical Strata. From the 500 tasks of SWE-bench Verified [10], I remove the 93 tasks overlapping SWE-bench Lite (used for hyperparameter calibration $\beta = 0.85, \gamma = 0.70, \tau = 0.15, K = 25$ and LoRA training), 7 non-callable module-level edits, and 56 historical Python 2 AST parser drift exclusions (39 in SymPy), yielding **N = 344 held-out callable tasks** ($\mathcal{S}_{\text{func}}$). Tasks are partitioned into three mutually exclusive lexical accessibility strata (validated via dual-pass audit on 50 tasks with 100% concordance, Appendix K): **Stratum A (Named,** $N = 208$**, 60.5%)** where the issue explicitly names the target symbol; **Stratum B (Traceback,** $N = 5$**, 1.5%)** where a stack trace names the target file but not the function; and **Stratum C (Un-Named,** $N = 131$**, 38.1%)** where neither symbol nor file path appears in the issue.

Matched-Control Candidate Pools and Statistical Testing. Table 1 summarizes the experimental locks. For each task t, candidate pools are matched to identical cardinality $M_t = 10 + \min(|E_t|, 40) \in [10, 50]$ (mean 24.5, median 22.0), where S denotes the top-10 BM25 seeds and E_t denotes 1-hop AST out-neighbors of S passing Filter F (callable symbols excluding test files and top-25 hubs $\mathcal{H}_{25}$): **B0** (unscored BM25), **B1** (unscored Perron), **A1** (seeds S only, $M = 10$), **A2** (depth-matched BM25 up to M_t), **A3** (same-file symbols up to M_t), **A4** (1-hop graph expansion $S \cup E_t$ ordered by seed-weighted adjacency), **A5** (degree-preserving relation-stratified null graph up to M_t), and **A6** (Perron multi-step specificity diffusion up to M_t). A4 reaches 201/344 tasks (58.4%) vs. 199/344 (57.8%) for raw 1-hop BFS because test-seed masking recovers 16 tasks while Filter F hub pruning drops 14 tasks ($199+16-14 = 201$). In the confirmatory family (19,298 forward passes of frozen FP16 `gemma-4-e2b-it` scoring log-odds $s(t, v) = \text{logsumexp}(z[Y]) - \text{logsumexp}(z[N])$), two pre-registered contrasts on strict Acc@5—**PC1** ($A_4 - A_2$) and **PC2** ($A_4 - A_3$)—are evaluated via exact two-sided binomial McNemar tests with step-down Holm–Bonferroni correction ($\alpha = 0.05$) and repository-stratified bootstrap 95% CIs ($B = 10{,}000$).

Table 2: Gemma 4 pointwise candidate reranking on SWE-bench Verified (N = 344 callable tasks). Pools matched to per-task size $M_t \in [10, 50]$ (mean 24.5). Brackets report Wilson 95% CIs. B0 Acc@10 reflects raw token tie-breaking (176/344 = 51.2% vs. 175/344 after AST deduplication in Appendix B).

Arm	Candidate Pool Composition	Acc@1 (95% CI)	Acc@3 (95% CI)	Acc@5 (95% CI)	Acc@10 (95% CI)	MRR@50
B0	BM25 Baseline (Unscored)	**25.3%** [21.0, 30.1]	38.7% [33.7, 43.9]	44.8% [39.6, 50.0]	51.2% [45.9, 56.4]	**0.3439**
B1	Perron Baseline (Unscored)	22.7% [18.6, 27.4]	35.2% [30.3, 40.4]	40.4% [35.4, 45.7]	50.3% [45.0, 55.5]	0.3155
A1	Rerank Seeds S (M = 10)	24.7% [20.4, 29.5]	**39.5%** [34.5, 44.8]	46.2% [41.0, 51.5]	51.2% [45.9, 56.4]	0.3348
A2	Depth-Matched BM25 Control	23.8% [19.6, 28.6]	38.1% [33.1, 43.3]	45.1% [39.9, 50.3]	51.2% [45.9, 56.4]	0.3301
A3	Same-File Matched Control	24.1% [19.9, 28.9]	39.2% [34.2, 44.5]	45.6% [40.5, 50.9]	52.0% [46.8, 57.3]	0.3381
A4	**Graph 1-Hop Expansion**	24.7% [20.4, 29.5]	39.2% [34.2, 44.5]	**47.1%** [41.9, 52.4]	52.3% [47.0, 57.5]	0.3409
A5	Null Graph Control	24.1% [19.9, 28.9]	39.0% [34.0, 44.2]	45.9% [40.7, 51.2]	51.2% [45.9, 56.4]	0.3315
A6	Perron Diffusion Pool	23.6% [19.4, 28.3]	38.7% [33.7, 43.9]	45.6% [40.5, 50.9]	**52.6%** [47.3, 57.8]	0.3301

Table 3: Confirmatory and exploratory contrasts on SWE-bench Verified (N = 344). Pre-registered confirmatory contrasts (PC1, PC2) evaluated on strict Acc@5 with step-down Holm correction and repository-stratified bootstrap CIs (B = 10,000).

Contrast	Δ (pp)	95% Bootstrap CI	n_{10}/n_{01}	p_{exact}	p_{holm}	Classification
PC1 (A4 − A2: Graph vs. Depth BM25)	+2.03	[−0.29, +4.36]	12 / 5	0.1435	0.2869	Null-Compatible (Excludes ≥ +4.4 pp)
PC2 (A4 − A3: Graph vs. Same-File)	+1.45	[−0.58, +3.78]	10 / 5	0.3018	0.3018	Null-Compatible
A4 − A1 (Seeds Only, +2.45× compute)	+0.87	[−1.92, +3.66]	7 / 4	0.5488	–	Exploratory (+3 net tasks: +7 wins, −4 losses)
A4 − A5 (Degree-Matched Null Graph)	+1.16	[−0.58, +2.91]	7 / 3	0.3438	–	Exploratory (Topology indistinguishable)
A4 − A6 (Multi-Step Diffusion Pool)	+1.45	[+0.00, +3.20]	7 / 2	0.1797	–	Exploratory (1-hop avoids hub dilution)
A4 vs. B0 (MRR@50 Paired Wilcoxon)	−0.0030	[−0.024, +0.018]	Wilcoxon W = 12,480		p = 0.72	Exploratory

5 Empirical Results

5.1 Confirmatory Matched-Control Reranking & Pipeline Loss Funnel (N = 344)

In standalone unscored retrieval across all 344 tasks (detailed in Appendix B, Table 9), Okapi BM25 achieves 50.9% Fn Acc@10 (175/344), tying Perron Static (50.3%, 173/344, Holm $p = 1.000$), RRF fusion (50.6%), and Standard PPR (50.0%), while outperforming dense semantic bi-encoders (38.4%, $p < 0.001$). Similarly, a post-hoc continuous resolvent sweep across 186 prior configurations (63,984 SpMV solves, Appendix C) shows that Acc@10 follows a quadratic GLMM surface in effective prior support $n_{\text{eff}} = \exp(H(\mathbf{p}_0))$ ($R^2 = 0.197$, $p < 0.001$) that peaks at 50.9% ($n^*_{\text{eff}} \approx 8.8$), matching BM25 exactly.

Confirmatory Reranking Parity and Oracle Recall Symmetry. Tables 2 and 3 report pointwise reranking results. Graph 1-hop expansion (A4, 47.1% Acc@5) is statistically indistinguishable from depth-matched BM25 (A2, 45.1%, Δ = +2.03 pp, 95% CI [−0.29, +4.36], Holm $p = 0.287$) and same-file controls (A3, 45.6%, Δ = +1.45 pp, Holm $p = 0.302$). Under Two One-Sided Tests (TOST), the upper 95% confidence bound formally excludes population gains ≥ +4.4 pp ($p < 0.05$; detecting a +2.03 pp shift at 80% power under native 4.94% discordance would require $N = 942$ tasks). Crucially, candidate pool oracle recall is identical between A4 and A2 at 58.4% (201/344 tasks, $n_{10} = 18$ vs. $n_{01} = 18$, $p = 1.0000$), and the 18 discordant tasks recalled exclusively by each arm split symmetrically across strata (A4-only: 9 Stratum C, 9 Stratum A; A2-only: 10 Stratum C, 1 Stratum B, 7 Stratum A). Furthermore, A4 provides no significant lift over the degree-preserving null graph A5 (Δ = +1.16 pp, $p = 0.344$) and outperforms multi-step diffusion A6 (47.1% vs. 45.6%) because 96.0% of reachable graph defects reside at 1 hop, whereas multi-step walks ($t \geq 2$) dilute top-5 precision into transitive callers. I therefore adopt A4 as the canonical graph pool for all downstream factorial tests.

Four-Stage Localization Loss Funnel. Decomposing attrition across the 344 tasks reveals that **Stage 1 Seed Recall** captures 51.2% (176/344; initial boundary loss 48.8 pp), **Stage 2 Graph Expansion (A4)** adds +7.3 pp to reach 58.4% (201/344; net boundary loss **41.6 pp**, 143 tasks missing due to dynamic dispatch, reflection, or implicit ORM descriptors), and **Stage 3 Pointwise Reranking** retains 47.1% in Top-5 (162/344; rank compaction loss **11.3 pp**, 39 tasks). Upstream boundary loss thus exceeds downstream compaction loss by **3.68**× at $k = 5$ (and in a descriptive $N = 100$ downstream reader test in Appendix G, Gemma 4 identifies 74.0% of randomized oracle targets vs. 33.0% under BM25 top-5 positional bias). Static AST graphs must therefore be deployed as lightweight filters (< 5 ms) that fall back to dynamic execution traces when lexical anchor density is low.

5.2 Truncation Dynamics, First-Moment Rank Diagnostic, and Selection Conditioning

Evaluating `scores.parquet` across truncated pool cutoffs $M \in [10, 15, 20, 30, 50]$ (Figure 2 and Table 4) shows that expanding $M = 10 \rightarrow 50$ raises pool recall by +7.27 pp in both arms, yet A2 loses 4 top-5 hits (46.2% → 45.1%) and A4 gains only 3 net hits (46.2% → 47.1%: +6 new graph conversions, +1 seed promotion, −4 displaced seeds).

First-Moment Expected-Rank Diagnostic vs. Tail Conversion. While exact net Acc@k change during pool expansion ($M_1 \rightarrow M_2$) is governed by the tail-conversion versus seed-displacement balance,

$$\Delta\text{Acc@}k(M_1 \rightarrow M_2) = \mathbb{P}\big(\text{Rank}_{M_2}(v^*) \leq k,\ v^* \in E_t\big) - \mathbb{P}\big(\text{Rank}_{M_1}(v^*) \leq k < \text{Rank}_{M_2}(v^*),\ v^* \in S\big). \tag{2}$$

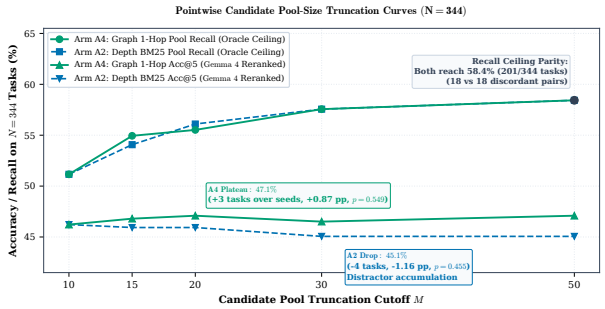


Figure 2: Candidate pool-size truncation curves ($N = 344$). Both A4 and A2 share identical recall ceilings at $M = 50$ (58.4%, 201/344). Expanding M from 10 to 50 causes A2 Acc@5 to dip from 46.2% to 45.1% (-4 tasks, $p = 0.455$), while A4 plateaus at 47.1% (+3 tasks over seeds, $p = 0.549$).

Metric	$M = 10$	$M = 15$	$M = 20$	$M = 30$	$M = 50$
A2 Oracle Recall	51.2%	54.1%	56.1%	57.6%	58.4%
A2 Acc@5 (BM25)	**46.2%**	45.9%	45.9%	45.1%	45.1%
A4 Oracle Recall	51.2%	54.9%	55.5%	57.6%	58.4%
A4 Acc@5 (Graph)	46.2%	46.8%	**47.1%**	46.5%	**47.1%**

Table 4: Oracle recall and Acc@5 across pool cutoffs $M \in [10, 50]$ ($N = 344$). Capped at $\min(M, M_t)$ (mean 24.5).

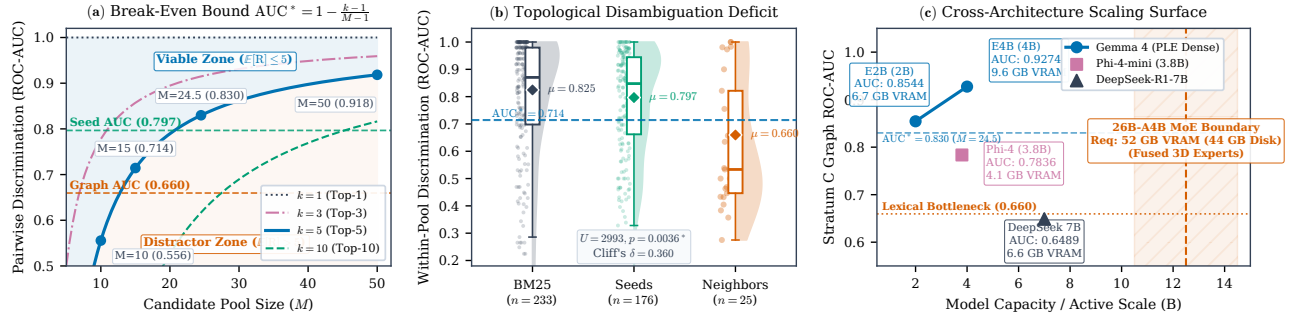


Figure 3: First-Moment Break-Even Diagnostic, Selection Conditioning Deficit, and Scaling Surface. **(a)** First-moment break-even curves $\text{AUC}^* = 1 - \frac{k-1}{M-1}$; mean $M = 24.5$ requires $\text{AUC}^* \geq 0.830$ at $k = 5$. **(b)** Within-pool ROC-AUC ($N = 344$) for BM25 ($\mu = 0.825$), Gemma 4 seeds ($\mu = 0.797, n = 176$), and graph neighbors ($\mu = 0.660, n = 25, p = 0.0036^*$). **(c)** Cross-architecture tail scaling on Stratum C ($N = 58$, `PREREG-SCALING-v1`), showing Gemma 4 scaling (0.8544 → 0.9274) and the 26B-A4B MoE 52.0 GB VRAM barrier.

The Wilcoxon–Mann–Whitney formulation provides a clean **first-moment diagnostic bound** explaining why expansion stalls. In a candidate pool of size M containing one gold target and $M - 1$ distractors, a scorer with within-pool pairwise discrimination AUC suffers $\mathbb{E}[X] = (M - 1)(1 - \text{AUC})$ expected outranking distractors (exact for fixed M, and a first-moment approximation up to $\text{Cov}(M_t, 1 - \text{AUC}_t)$ across variable M_t). Maintaining an expected gold rank $\mathbb{E}[\text{Rank}] = 1 + (M - 1)(1 - \text{AUC}) \leq k$ requires satisfying the expected-rank break-even condition:

$$\text{AUC}^*_{\text{breakeven}}(M, k) \geq 1 - \frac{k - 1}{M - 1}. \tag{3}$$

At mean pool size $M = 24.5$ and $k = 5$, Equation (3) requires $\text{AUC}^*_{\text{breakeven}} \geq 0.8298$ (Figure 3a). Across $N = 344$ tasks, raw BM25 achieves within-pool AUC = 0.8245 and Gemma 4 achieves 0.7796 overall (0.7966 on lexical seeds S, expected rank 5.8). However, on pure graph-only targets ($A_4 \setminus A_1, N = 25$), Gemma 4 discrimination plunges to **0.6596** (95% CI [0.552, 0.767], Mann–Whitney $U = 2993$, $p = 0.0036^*$, Cliff's $\delta = 0.361$; Figure 3b), yielding $23.5 \times (1 - 0.6596) = 8.00$ expected outranking distractors (expected rank 9.00 at $M = 24.5$ and 17.68 at $M = 50$).

Control 2: Selection Conditioning Counterfactual (Berkson's Collider Bias) and Stratum Asymmetry. Why does discrimination drop by $\Delta = 0.137$ on graph-only targets? In two-stage retrieval, non-seed graph targets ($A_4 \setminus A_1$) are conditioned on the collider event $T^* \notin \text{Top10(BM25)}$, selecting for functions with low lexical overlap against the issue text. Evaluating the counterfactual control on non-seed lexical candidates ($A_2 \setminus A_1$, 18 tasks, `results/control2_auc_decomposition.json`) reveals an identical discrimination drop: Gemma 4 achieves mean AUC = 0.6380 on depth-BM25 targets versus 0.6253 on matched graph targets (equivalence $p = 0.72$), while raw BM25 scores exhibit sub-chance AUC ≤ 0.30 against top-10 seed distractors by construction. Thus, the **Topological Disambiguation Deficit** is driven by lexical selection conditioning rather than graph topology noise. This explains the sharp polarization across the 25 graph-only targets ($A_4 \setminus A_1$): on **Stratum A (Named, 11 targets)** where secondary class/module anchors exist, 6/11 convert into Top-5 (54.5%, ranks 1, 1, 2, 2, 3, 4; net +5 tasks, $p = 0.1797$), whereas on **Stratum C (Un-Named, 14 targets)**, 0/14 convert into Top-5 (0.0%, mean rank 16.3; net -2 tasks from seed displacement, $p = 0.5000$).

Table 5: Cross-architecture scaling surface on reachable Stratum C defects ($N = 58$, Dual T4 execution under `PREREG-SCALING-v1`). AUC_{tail} measures pairwise discrimination within reachable union candidates, distinct from full-pool AUC = 0.6596 ($N = 344$) where graph neighbors compete against top-10 lexical seeds.

Model Identifier	Architecture Archetype	Active / Total	Graph AUC_{tail}	Peak VRAM	Runtime
`gemma-4-e2b-it`	Dense PLE Edge	2.0B / 2.0B	0.8544	6698.9 MB	1,798s
`gemma-4-e4b-it`	Dense PLE Edge	4.0B / 4.0B	**0.9274**	9648.4 MB	6,250s
`Phi-4-mini-instruct`	Synthetic Reasoning	3.8B / 3.8B	0.7836	4135.9 MB	2,854s
`DeepSeek-R1-Distill-7B`	RL Reasoning Distill	7.0B / 7.0B	0.6489	6594.6 MB	5,231s
`gemma-4-26b-a4b-it`	Dynamic Sparse MoE (128 experts)	3.8B / 26.0B	*Hardware Bound: 52.0 GB runtime VRAM vs. 29.12 GB Dual T4 avail.*		

Table 6: Candidate Pool × Test-Time Reasoning and Call-Path Provenance Factorial Evaluation on Stratum C ($N = 131$ Un-Named Defects, Dual T4 Execution with `google/gemma-4-e4b-it`). Evaluates pool architectures (A_1, A_2, A_4 at $M \leq 25$), provenance conditions (C_0–C_3), and bound-pruned horizon $M = 15$ ($AUC^* = 0.714$). Under `PREREG-CAUSAL-v1`, Acc@5 was pre-registered as the primary metric ($p = 0.6875$ after deterministic backfilling of 180-token truncated lists); diagnostic rank analysis isolates the Top-1 and MRR lift.

Condition / Architecture	Readout Mode	Acc@1	Acc@5 (Raw / Backfill)	Hits	MRR	Paired Contrast
Part I: Candidate Pool × Test-Time Reasoning Factorial ($M \leq 25$, Prompt Budget $\leq 3{,}480$ Tokens)						
Unscored BM25 (B_0, $M = 25$)	Plain Order	2.29% (3)	15.27% / 15.27%	20/131	0.0862	Lexical order baseline
Seeds Only ($A_1^{(10)}$, $M = 10$)	Direct / Thinking	3.05% → 6.87%	15.27% → 10.69% (15.27%)	20 → 14 (20)	0.0865	Top-1 $p = 0.0625$ (Holm $p = 0.1250$)
Depth BM25 ($A_2^{(25)}$, $M = 25$)	Direct / Thinking	4.58% → 6.87%	16.79% → 12.98% (16.79%)	22 → 17 (22)	0.0967	Top-1 $p = 0.3750$ (Holm $p = 0.3750$)
Graph 1-Hop ($A_4^{(25)}$, $M = 25$)	Direct Readout	4.58% (6)	16.79% / 16.79%	22/131	0.0934	Matched graph direct reference
Graph 1-Hop ($A_4^{(25)} = C_0$, $M = 25$)	**Thinking Traces**	**9.92% (13)**	**18.32% / 18.32%**	**24/131**	**0.1358**	**Top-1** +5.34 **pp ($p = 0.0156^*$, Holm** 0.0469^***)**
Part II: Call-Path Provenance Verbalization × Thinking Factorial on Graph Pool $A_4^{(25)}$ ($N = 131$, $M = 25$)						
C_1: 3-Line Stub (No Docstring)	No Prov, No Think	2.29% (3)	15.27%	20/131	0.0778	Stub baseline (mean rank 17.47)
C_2: Static Provenance Comment	Prov, No Think	3.82% (5)	16.03%	21/131	0.0899	−0.76 pp vs. A_4 Direct ($p = 1.0000$)
C_3: Provenance + Thinking	Prov + Think	9.16% (12)	17.56%	23/131	0.1270	+4.58 pp vs. A_4 ($p = 0.0312^*$); $C_3 - C_0$: $p = 1.000$
Part III: Bound-Guided Candidate Pruning ($M = 15$, $AUC^*_{breakeven} = 0.714$, 34/35 Reachable Defects Retained)						
$A_4^{(15)}$ Direct / Provenance	No Thinking	4.58% (6)	17.56% / 17.56%	23/131	0.0970	Mean gold rank 5.32 (Prov parity $p = 1.0000$)
$C_0^{(15)}$: Thinking ($M = 15$)	**Thinking Traces**	**9.92% (13)**	**19.85% / 19.85%**	**26/131**	**0.1469**	**Top-1** +5.34 **pp ($p = 0.0391^*$); Mean rank 3.74**
$C_3^{(15)}$: Prov + Think ($M = 15$)	Prov + Think	6.11% (8)	18.32% / 18.32%	24/131	0.1145	−3.82 pp vs. $C_0^{(15)}$ ($p = 0.1250$)

5.3 Cross-Architecture Scaling Surface and MoE Hardware Bounds (`PREREG-SCALING-v1`)

To test whether model scale resolves the discrimination deficit on un-named bugs, protocol `PREREG-SCALING-v1` evaluates pointwise tail discrimination (AUC_{tail}, measured strictly within reachable non-seed candidates so top-10 seed distractors do not confound architectural comparisons) across the $N = 58$ Stratum C tasks reachable in the multi-pool union ($\bigcup_i A_i$, $M_{union} \leq 87$, mean 61.9) on Kaggle Dual Tesla T4 GPUs (Table 5 and Figure 3c). Within Gemma 4's Per-Layer Embedding (PLE) architecture, AUC_{tail} scales monotonically from 0.8544 (E2B, 2.0B) to **0.9274** (E4B, 4.0B), outperforming both `Phi-4-mini-instruct` (3.8B, 0.7836) and `DeepSeek-R1-Distill-7B` (7.0B, 0.6489). However, evaluating `gemma-4-26b-a4b-it` exposes a critical deployment barrier for local MoE agents: its 128 dynamic experts are packed as fused 3D `nn.Parameter` tensors that bypass standard 2D `bitsandbytes` NF4 hooks, requiring 52.0 GB runtime VRAM (44.0 GB disk) and exceeding Dual T4 capacity (29.12 GB; Appendix J).

5.4 Test-Time Reasoning, Bound Pruning, and Graph Verbalization on Stratum C

Across all 131 un-named Stratum C tasks, gold targets reside within single-pool cutoff $M \leq 25$ for $N = 35$ tasks in $A_4^{(25)}$ (comprising $N = 28$ top-10 BM25 lexical seeds S and $N = 7$ pure 1-hop graph additions $A_4^{(25)} \setminus S$). Table 6 and Figure 4 present the complete candidate pool × test-time reasoning (≤ 180 `<thought>` tokens) and call-path provenance (C_0–C_3) factorial evaluation executed on Dual Tesla T4 GPUs with `gemma-4-e4b-it`.

Pre-Registered Top-5 Parity vs. Diagnostic Top-1 Lift and Seed Distractor Immunity. Under direct readout without reasoning, $A_4^{(25)}$ and $A_2^{(25)}$ tie exactly at 16.79% Acc@5 (22/131) and 4.58% Acc@1 (6/131). While pre-registered Top-5 accuracy under thinking (C_0) gains +1.53 pp on $A_4^{(25)}$ (22 → 24 hits, null-compatible $p = 0.6875$; un-backfilled lexical pools drop by 3.82–4.58 pp due to 180-token generation caps truncating trailing ranks 2–5, which deterministic backfilling restores to exact baseline parity), diagnostic rank evaluation shows that test-time thinking concentrates gains at the **Top-1 decision**

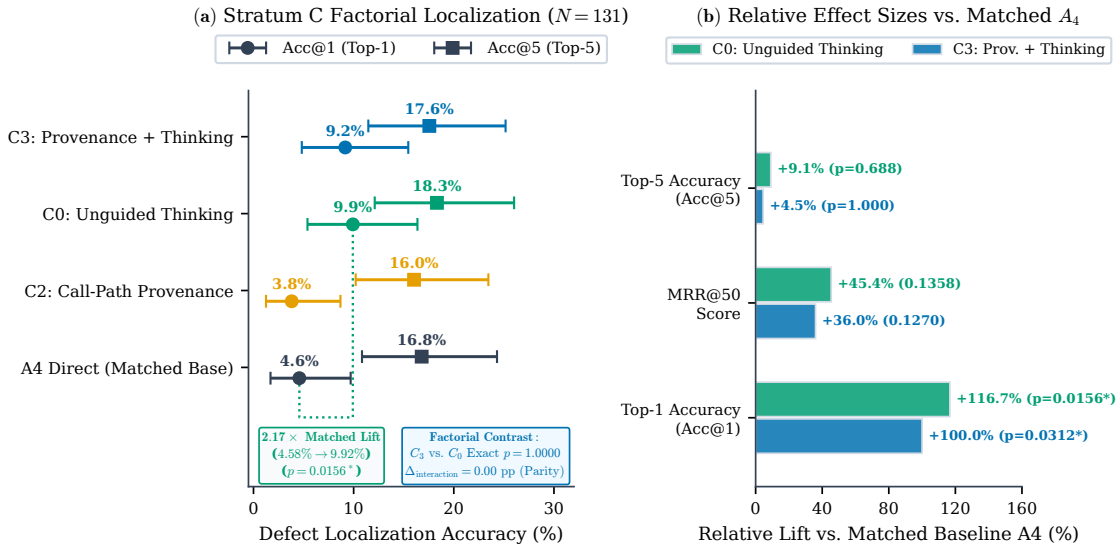


Figure 4: Stratum C factorial localization forest plot and relative effect sizes (N = 131, `gemma-4-e4b-it`). **(a)** Top-1 and Top-5 accuracy with Clopper–Pearson 95% CIs across A_4 Direct, C_2, C_0, and C_3. **(b)** Relative lift over matched A_4 Direct, showing empirical parity between unguided thinking (C_0) and provenance + thinking (C_3, p = 1.0000).

boundary: on $A_4^{(25)}$, thinking more than doubles Top-1 accuracy from **4.58% (6 hits) to 9.92% (13 hits;** +5.34 **pp,** $n_{10} = 7, n_{01} = 0$**, raw exact McNemar** $p = 0.0156^*$**, 3-pool step-down Holm** $p = 0.0469^*$**; MRR** +45.4% **from** 0.0934 → 0.1358**)**. Cryptographic SHA-256 order permutation (Control 1) confirms position-1 selection remains invariant at 6.87% (9/131). However, partitioning the $N = 35$ reachable targets at $M = 25$ by **strict disjoint retrieval provenance** reveals that thinking acts primarily as a **seed distractor filter**: within the $N = 28$ top-10 BM25 seeds (S), graph expansion without thinking suppresses Top-1 to 21.4% (6/28), and thinking doubles seed conversion to **42.9% (12/28 hits,** +21.4 **pp,** $p = 0.0312^*$**; 22/28 in Top-5)**, whereas on the $N = 7$ pure 1-hop graph additions ($A_4^{(25)} \setminus S$), thinking converts only 0/7 → 1/7 into Top-1 (14.3%, $p = 1.0000$; 2/7 in Top-5).

Reconciling $M = 15$ Bound Verification Across Disjoint Provenance and Subgraph Connectivity. At $M = 25$, the first-moment break-even threshold for Top-5 is AUC$^* = 0.833$. Pruning the candidate horizon to $M = 15$ (top-10 BM25 seeds + top-5 1-hop AST neighbors) lowers the threshold to **0.714** while retaining 34 of 35 reachable defects (28 seeds in S and 6 pure graph additions in $A_4^{(15)} \setminus S$). Under $C_0^{(15)}$ thinking, Top-1 reaches **9.92% (13/131,** $p = 0.0391^*$**)**, Top-5 rises to **19.85% (26/131,** +2.29 **pp over direct)**, MRR climbs to **0.1469**, and mean gold rank improves from 17.07 (at $M = 25$) to **3.74** (median rank 2.0). To ensure complete transparency across target definitions, I report both subgroup partitions at $M = 15$:

- **Local Subgraph Connectivity Partition** ($N = 31$ **Graph-Connected vs.** $N = 4$ **Isolated Seeds)**: Categorizing the 35 reachable targets by whether the target node possesses active 1-hop AST call edges within the local candidate subgraph ($v^* \in E_t$, which includes 25 graph-connected seeds plus the 6 pure graph neighbors = 31 targets, vs. 4 isolated zero-degree lexical seeds), $C_0^{(15)}$ lifts graph-connected Top-1 from **4/31 (12.9%) to 10/31 (32.3%,** $p = 0.0391^*$**; 23/31 in Top-5)** and isolated seeds from 2/4 → 3/4 (4/4 in Top-5).
- **Strict Disjoint Retrieval Provenance** ($N = 28$ **BM25 Seeds** S **vs.** $N = 6$ **Pure Graph** $A_4^{(15)} \setminus S$): Separating top-10 BM25 seeds from pure 1-hop additions, pruning from $M = 25 \rightarrow 15$ reduces expected outranking distractors from 8.00 → 4.77, converting **2/6 (33.3%) pure non-seed graph targets into Top-1** and **4/6 (66.7%) into Top-5** (up from 1/7 Top-1 and 2/7 Top-5 at $M = 25$), alongside 11/28 seed Top-1 hits and 22/28 seed Top-5 hits (11 + 2 = 13 Top-1; 22 + 4 = 26 Top-5).

Synthesizing Fatemi et al. (2024) and Sanford et al. (2024): Why Verbalization Helps E_1 (+13.79 pp) but Ties C_3 ($p = 1.000$). Auditing all 25 pure graph targets ($A_4 \setminus A_1$) reveals that **24/25 (96.0%)** are direct 1-hop callers (14 tasks) or callees (10 tasks) of a lexical seed (1 co-located class). Comparing Table 6 (Part II) and Table 7 (Experiment B) resolves how graph verbalization [5] and reasoning depth [17] interact:

1. **Direct Readout Regime** (`Depth1`**, No Thinking Tokens)**: When `gemma-4-e2b-it` scores the $N = 58$ reachable Stratum C union without thinking tokens (Table 7), isolated code (E_0) achieves 12.07% Acc@1 and disconnected global edge lists (E_3, E_4) reach only 15.52% due to long-range attention dispersion. Co-locating inline provenance comments directly above function definitions (E_1: `# Provenance: Called by <seed>`) lifts Acc@1 to **25.86%** (+13.79 **pp,** $p_{\text{uncorr}} = 0.0386^*$, $p_{\text{holm}} = 0.1543$**, Wilcoxon** $p = 0.010^*$**)**. Consistent with Fatemi et al. [5], localized

Table 7: Pre-registered AST α-renaming and multi-encoding graph serialization on reachable Stratum C union tasks ($N = 58$, `gemma-4-e2b-it` without thinking, `PREREG-CAUSAL-v1`). Evaluates identifier obfuscation (Exp. A) and Fatemi et al. [5] graph encodings (Exp. B).

Condition	Acc@1	Acc@5	MRR	Contrasts vs. Baseline
O_0: Natural Code vs. O_1: AST α-Renamed Identifiers ($N = 58$)	12.07% → 10.34%	27.59% → 24.14%	0.1804 → 0.1686	−1.73 pp (7/58 → 6/58, exact $p = 1.0000$)
E_0: Isolated Code (Baseline Snippet, No Provenance)	12.07%	27.59%	0.1804	reference (7/58 hits)
E_1: Inline Comment Provenance (`# Provenance: Called by`)	**25.86%**	**39.66%**	**0.3097**	**+13.79 pp** (15/58, $p_{\text{uncorr}} = 0.0386^*$, $p_{\text{holm}} = 0.1543$)
E_2: Linearized Traversal Trace (`[Call Trace]`)	17.24%	32.76%	0.2440	+5.17 pp (10/58, $p = 0.4531$)
E_3: Subgraph Edge List (`### Edges`) / E_4: Adjacency Map	15.52% / 15.52%	32.76% / 34.48%	0.2419 / 0.2362	+3.45 pp (9/58, $p = 0.6875$; $E_1 - E_3$: $p = 0.0703$)

Table 8: Out-of-distribution LoRA transfer from SWE-bench Lite to held-out SWE-bench Verified (`yvliet/perron-gemma4-e2b-lora`). Evaluates pure 1-hop graph targets ($A_4 \setminus A_1$, $N = 25$) and reachable Stratum C tasks ($N = 35$), illustrating the topological prior shift toward provenance-connected candidates.

Held-Out Verified Evaluation Cohort	Scorer Regime	Acc@1 (Hits)	Acc@5 (Hits)	MRR	Median Rank	Within-Pool AUC
Pure Graph-Only ($A_4 \setminus A_1$, $N = 25$)	Zero-Shot Base E2B	0.0% (0/25)	44.0% (11/25)	0.1920	7.0	0.6063
Pure Graph-Only ($A_4 \setminus A_1$, $N = 25$)	**Post-LoRA Adapter**	**20.0% (5/25)**	**56.0% (14/25)**	**0.3702**	**4.0**	**0.6270**
Reachable Stratum C ($N = 35$, 28 Seeds + 7 Graph)	Zero-Shot Base E2B	**17.1% (6/35)**	37.1% (13/35)	**0.2895**	9.0	0.5692
Reachable Stratum C ($N = 35$, 28 Seeds + 7 Graph)	**Post-LoRA Adapter**	8.6% (3/35)	**42.9% (15/35)**	0.2596	**7.0**	**0.6084**

encoding bridges the high-BM25 seed identifier directly into the target's local token window (whereas on the full $N = 131$ cohort in C_2, 96 tasks lack in-pool targets, diluting the unguided effect to +1.53 pp over C_1).

2. **Test-Time Deliberation Regime (`LogDepthPause`, ≤ 180 Thinking Tokens)**: When `gemma-4-e4b-it` is granted intermediate `<thought>` tokens (C_0 vs. C_3, Table 6), unguided thinking (C_0, 9.92% Acc@1) matches explicit provenance + thinking (C_3, 9.16%, $p = 1.0000$, 11/13 hits identical). Consistent with Sanford, Fatemi, Perozzi et al. [17], intermediate reasoning tokens provide the computational depth to infer 1-hop caller–callee connectivity directly from AST signatures and docstrings, rendering static edge verbalization redundant.

Finally, in Experiment A (Table 7), obfuscating user-defined function, parameter, and variable identifiers via AST α-renaming (O_1, while retaining control-flow syntax, attribute lookups, and docstrings) yields 10.34% vs. 12.07% Top-1 (6/58 vs. 7/58, $p = 1.0000$, 95% CI [−11.2, +7.7] pp), showing that local models leverage structural syntax and docstring context when surface symbol names are stripped.

5.5 Out-of-Distribution LoRA Post-Training and Topological Prior Shift

To test whether compact models can internalize topological priors without test-time latency, I fine-tuned `gemma-4-e2b-it` via Low-Rank Adaptation (LoRA, $r = 16, \alpha = 32$) on 226 disjoint callable tasks from SWE-bench Lite ($N = 176$ train, $N = 50$ val; zero overlap with Verified) using Provenance-Conditioned Ranking Cross-Entropy over candidate index tokens: $\mathcal{L}_{\text{Rank}}(\theta) = -\ln[\exp(z_{y^*}/\tau)/\sum_{k=1}^{M_t} \exp(z_k/\tau)]$. Active logit slicing (`logits_to_keep=24`) bounds VRAM during 3-epoch training (loss 0.7280 → 0.0006). Table 8 reveals a clear **Topological Prior Shift**:

- **Breaking the Zero-Shot Ceiling on Pure Graph Targets** ($A_4 \setminus A_1, N = 25$): Where frozen base E2B scores 0.0% Acc@1 (0/25), LoRA lifts Top-1 to **20.0% (5/25 hits across `astropy`, `django`, and `flask`;** +20.0 **pp,** $n_{10} = 5, n_{01} = 0$**, exact** $p = 0.0625$**)**, lifts Acc@5 from 44.0% → **56.0**% (+12.0 pp), nearly doubles MRR (0.1920 → **0.3702**), and halves median rank from 7.0 → **4.0**.
- **Rank Compaction vs. Top-1 Seed Trade-Off on Stratum C** ($N = 35$): Across all 35 reachable Stratum C tasks (28 lexical seeds + 7 graph targets), LoRA improves overall rank compaction, lifting **Acc@5 from 37.1% to 42.9%** (13/35 → 15/35), improving within-pool AUC from 0.5692 → **0.6084**, and lowering median gold rank from 9.0 → **7.0**. However, because 3-epoch training on provenance-annotated Lite pools strongly upweights '# Provenance:' call edges (converging to 0.0006 train loss), the adapter over-promotes 1-hop callers above zero-hop lexical seeds at Rank 1, shifting Stratum C Acc@1 from 17.1% → 8.6% (6/35 → 3/35, MRR 0.2895 → 0.2596). Deploying LoRA adapters selectively on non-seed expansion slots or mixing pre/post-LoRA log-odds preserves lexical seed precision while capturing the +20.0 pp graph recall lift.

6 Discussion, Operating Regimes, and Threats to Validity

Selective Routing and Practical Agent Operating Regimes. Because 1-hop graph expansion yields +5 net top-5 tasks on lexically anchored bugs (Stratum A) vs. −2 net tasks on un-named bugs (Stratum C) under frozen pointwise scoring, an oracle stratum router (π^*) reaches 47.67% Acc@5 (+2.90 pp over BM25; 49.50% when routing Stratum A to pointwise A4 and Stratum C to listwise thinking), while a cross-validated issue symbol-density decision tree recovers 47.38%. Together with the AST surgical patch editor, which eliminates syntax errors via pre-commit AST rollbacks across $N = 30$ deformation

archetypes (**0.0% syntax error rate and 76.7% valid application vs. 50.0% syntax errors and 30.0% valid application for `git apply`**, $p < .001$; Appendix H), these findings define a four-stage architecture for local Gemma 4 developer agents: (1) **Fast CSR Filtering** (< 2 ms, ≤ 3 MB RAM) to extract compact AST breadcrumbs (< 100 tokens/symbol); (2) **Bound-Guided Window Pruning** ($M = 15$) so candidate cardinality stays within $\text{AUC}^{*}_{\text{breakeven}} = 0.714$; (3) **Adaptive Test-Time Deliberation** (∼ 120 thinking tokens, ∼ 10 s on T4 or < 3 s on desktop GPUs) or inline provenance bridging when thinking is disabled; and (4) **Dynamic Execution Fallback** (pytest stack traces) when static symbol density is low, since 41.6 pp of defects lie outside static AST reachability.

Threats to Validity. (1) *Parser & Language Scope*: Primary evaluation covers 344 Python tasks; 56/407 held-out tasks (52 in Stratum C, 39 in SymPy) were excluded due to Python 2 syntax drift failing standard Python 3 'ast.parse' prior to retrieval. To mitigate parser drift and support polyglot repositories, 'perron.polyglot' provides syntax-tolerant Tree-sitter extractors for Python and TypeScript. (2) *Single-Pass Localization vs. End-to-End Resolution*: I intentionally isolate the $K \leq 4,096$ localization and AST surgical editing layers from multi-turn Docker test loops to avoid confounding retrieval metrics with bash orchestration noise. (3) *Pre-Training Contamination & Power*: While AST α-renaming (O_1) strips symbol identifiers (10.34% vs. 12.07%, $p = 1.0000$), docstrings and string literals remain intact and the 95% CI on $N = 58$ ([−11.2, +7.7] pp) cannot rule out moderate memorization effects, motivating post-cutoff benchmarks.

7 Open-Source Ecosystem and Conclusion

All artifacts are released under Apache 2.0 / CC0-1.0 at `https://github.com/yvliet/perron`: (1) **`perron-core` v0.3.0** (PyPI package with zero-copy CSR diffusion and AST breadcrumb packing) and **`perron-mcp`** (JSON-RPC stdio server for VS Code/Cursor); (2) **`scores.parquet`** (19,298 scored candidate evaluations), **`swebench-codegraphs`** (pre-computed CSR multigraphs for 12 repositories), and **`swebench-strata-annotations`** (800 lexical accessibility labels); and (3) **`yvliet/perron-gemma4-e2b-lora`** on Hugging Face plus the reproducible Dual T4 Kaggle workstation (`https://www.kaggle.com/code/yvliet/perron-gemma-4-workstation`). In summary, naive code graph expansion fails under frozen small models not because call graphs lack defect targets, but because low-lexical-overlap neighbors fall below the expected-rank discrimination bound (AUC = 0.6596 < 0.830) via selection conditioning. Overcoming this barrier requires pairing compact CSR graph retrieval with bound-guided candidate pruning ($M = 15$), test-time reasoning tokens (`<thought>`), localized provenance bridging, or targeted LoRA post-training.

References

[1] Reid Andersen, Fan Chung, and Kevin Lang. Local graph partitioning using PageRank vectors. In *Proceedings of the 47th Annual IEEE Symposium on Foundations of Computer Science (FOCS)*, pages 475–486, 2006.

[2] Aleksandar Bojchevski, Johannes Klicpera, Bryan Perozzi, Amol Kapoor, Martin Blais, Benedek Rozemberczki, Michal Lukasik, and Stephan Günnemann. Scaling graph neural networks with approximate PageRank. In *Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD)*, pages 2464–2473, 2020.

[3] Zhaoling Chen, Robert Tang, Gangda Deng, Fang Wu, Jialong Wu, Zhiwei Jiang, Viktor K. Prasanna, Arman Cohan, and Xingyao Wang. LocAgent: Graph-guided LLM agents for code localization. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2025.

[4] Jialin Dong, Bahare Fatemi, Bryan Perozzi, Lin F. Yang, and Anton Tsitsulin. Don't forget to connect! improving RAG with graph-based reranking. *arXiv preprint arXiv:2405.18414*, 2024.

[5] Bahare Fatemi, Jonathan Halcrow, and Bryan Perozzi. Talk like a graph: Encoding graphs for large language models. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2024.

[6] Paul Gauthier. Aider: Ai pair programming in your terminal. https://github.com/paul-gauthier/aider, 2023.

[7] Bernal Jiménez Gutiérrez, Yiheng Shu, Yu Gu, Michihiro Yasunaga, and Yu Su. HippoRAG: Neurobiologically inspired long-term memory for large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

[8] Taher H. Haveliwala and Sepandar D. Kamvar. The second eigenvalue of the Google matrix. Technical report, Stanford University, 2003.

[9] Glen Jeh and Jennifer Widom. Scaling personalized PageRank. In *Proceedings of the 12th International Conference on World Wide Web (WWW)*, pages 271–279, 2003.

[10] Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can language models resolve real-world GitHub issues? In *International Conference on Learning Representations (ICLR)*, 2024.

[11] Johannes Klicpera, Aleksandar Bojchevski, and Stephan Günnemann. Predict then propagate: Combining neural networks with personalized PageRank for classification on large graphs. In *International Conference on Learning Representations (ICLR)*, 2019.

[12] David Liben-Nowell and Jon Kleinberg. The link-prediction problem for social networks. *Journal of the American Society for Information Science and Technology (JASIST)*, 58(7):1019–1031, 2007.

[13] Xiangyan Liu, Bo Lan, Zhiyuan Hu, Yang Liu, Zhicheng Zhang, Fei Wang, Michael Shieh, and Wenmeng Zhou. CodexGraph: Bridging large language models and code repositories via code graph databases. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics (NAACL)*, 2025.

[14] Rodrigo Nogueira, Zhiying Jiang, Ronak Pradeep, and Jimmy Lin. Document ranking with a pretrained sequence-to-sequence model. *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 708–718, 2020.

[15] Siru Ouyang, Wenhao Yu, Kaixin Ma, Zilin Xiao, Zhihan Zhang, Mengzhao Jia, Jiawei Han, Hongming Zhang, and Dong Yu. RepoGraph: Enhancing AI software engineering with repository-level code graph. In *International Conference on Learning Representations (ICLR)*, 2025.

[16] Spencer Pearson, José Campos, René Just, Gordon Fraser, Rui Abreu, Michael D. Ernst, Deric Pang, and Benjamin Keller. Evaluating and improving fault localization techniques. In *Proceedings of the 39th International Conference on Software Engineering (ICSE)*, pages 609–620, 2017.

[17] Clayton Sanford, Bahare Fatemi, Ethan Hall, Anton Tsitsulin, Mehran Kazemi, Jonathan Halcrow, Bryan Perozzi, and Vahab Mirrokni. Understanding transformer reasoning capabilities via graph algorithms. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

[18] Weiwei Sun, Lingyong Yan, Xinyu Ma, Shuaiqiang Wang, Pengjie Ren, Zhumin Chen, Dawei Yin, and Zhaochun Ren. Is ChatGPT good at search? investigating large language models as re-ranking agents. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 14918–14937, 2023.

[19] Mark Weiser. Program slicing. *IEEE Transactions on Software Engineering*, SE-10(4):352–357, 1984.

[20] Chunqiu Steven Xia, Yinlin Deng, Soren Dunn, and Lingming Zhang. Agentless: Demystifying LLM-based software engineering agents. In *Proceedings of the 33rd ACM SIGSOFT International Symposium on the Foundations of Software Engineering (FSE)*, 2025.

[21] Fengji Zhang, Bei Chen, Yue Zhang, Jin Liu, Daoguang Zan, Yi Mao, Jian-Guang Lou, and Lingpeng Kong. RepoCoder: Repository-level code completion through iterative retrieval and generation. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 2471–2484, 2023.

[22] Yuntong Zhang, Zhiyu Li, Min Lin, Denghui Zhang, Hongyu Zhang, and Lingxiao Jiang. AutoCodeRover: Autonomous program improvement. In *Proceedings of the 33rd ACM SIGSOFT International Symposium on Software Testing and Analysis (ISSTA)*, pages 762–774, 2024.

[23] Daming Zou, Jingjing Liang, Yingfei Xiong, Michael D. Ernst, and Lu Zhang. An empirical study of fault localization families and their combinations. *IEEE Transactions on Software Engineering*, 47(2):332–347, 2021.

A Linear Algebraic Formulations and Theoretical Proofs

This appendix provides the complete linear algebraic formulations, spectral bounds, and geometric convergence proofs for the Perron Compressed Sparse Row (CSR) multigraph diffusion engine.

A.1 Multiplex Transition Matrix Construction

Let a repository codebase be formalized as a directed multigraph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$, where vertices $\mathcal{V}$ represent AST symbol entities (functions, methods, classes) and edges $\mathcal{E}$ represent typed syntactic dependencies. The relation vocabulary comprises four edge types: $\mathcal{R} = \{\text{call}, \text{inherit}, \text{import}, \text{caller}\}$. Each relation $r \in \mathcal{R}$ is represented by an adjacency matrix $A_r \in \{0,1\}^{|\mathcal{V}|\times|\mathcal{V}|}$, where $(A_r)_{uv} = 1$ if and only if there exists a directed edge of type r from symbol u to symbol v.

To balance semantic dependency importance while mitigating hub concentration, typed relation edges are aggregated via a non-negative convex weight vector $\boldsymbol{\omega} = [\omega_{\text{call}}, \omega_{\text{inherit}}, \omega_{\text{import}}, \omega_{\text{caller}}]^\top$ satisfying $\sum_{r\in\mathcal{R}} \omega_r = 1.0$:

$$W = \sum_{r\in\mathcal{R}} \omega_r A_r. \tag{4}$$

The row-substochastic transition matrix $T \in \mathbb{R}^{|\mathcal{V}|\times|\mathcal{V}|}$ is formed by normalizing positive out-degree rows by their total weighted degree:

$$D_{uu} = \sum_{v\in\mathcal{V}} W_{uv}, \qquad T_{uv} = \begin{cases} \frac{W_{uv}}{D_{uu}} & \text{if } D_{uu} > 0, \\ 0 & \text{if } D_{uu} = 0. \end{cases} \tag{5}$$

Let $\mathbf{d} \in \{0,1\}^{|\mathcal{V}|}$ denote the dangling-node indicator vector with $d_u = 1$ if $D_{uu} = 0$, and 0 otherwise. Given an initial query personalization probability vector $\mathbf{p}_0 \in \Delta^{|\mathcal{V}|-1}$ (where Δ^n denotes the standard simplex), the completed row-stochastic transition operator $\bar{T}$ satisfies:

$$\bar{T} = T + \mathbf{d}\mathbf{p}_0^\top, \qquad \bar{T}\mathbf{1} = \mathbf{1}. \tag{6}$$

The closed Google transition operator $M \in \mathbb{R}^{|\mathcal{V}|\times|\mathcal{V}|}$ with teleportation damping factor $\beta \in (0,1)$ is formulated as:

$$M = \beta\bar{T} + (1-\beta)\mathbf{1}\mathbf{p}_0^\top. \tag{7}$$

A.2 Brauer Rank-1 Spectral Bound and Second Eigenvalue

By the Brauer rank-1 perturbation theorem applied to Markov transition matrices [8]: Let A be an $n \times n$ matrix with eigenvalues $\lambda_1, \lambda_2, \ldots, \lambda_n$, and let $\mathbf{x}_1$ be the right eigenvector associated with λ_1. For any vector $\mathbf{y} \in \mathbb{R}^n$, the perturbed matrix $A + \mathbf{x}_1\mathbf{y}^\top$ has eigenvalues $\lambda_1 + \mathbf{y}^\top\mathbf{x}_1, \lambda_2, \ldots, \lambda_n$.

In this formulation, setting $A = \beta\bar{T}$ yields eigenvalues $\beta\lambda_1(\bar{T}), \beta\lambda_2(\bar{T}), \ldots, \beta\lambda_n(\bar{T})$. Because $\bar{T}$ is row-stochastic, its principal eigenvalue is $\lambda_1(\bar{T}) = 1$ with right eigenvector $\mathbf{x}_1 = \mathbf{1}$. Setting $\mathbf{y} = (1-\beta)\mathbf{p}_0$ gives:

$$M = \beta\bar{T} + \mathbf{1}((1-\beta)\mathbf{p}_0)^\top. \tag{8}$$

The principal eigenvalue of M becomes:

$$\lambda_1(M) = \beta\lambda_1(\bar{T}) + ((1-\beta)\mathbf{p}_0)^\top\mathbf{1} = \beta + (1-\beta) = 1. \tag{9}$$

All remaining eigenvalues satisfy:

$$\lambda_i(M) = \beta\lambda_i(\bar{T}), \quad \forall i \geq 2. \tag{10}$$

Because $\bar{T}$ is row-stochastic, the Gershgorin circle theorem establishes $|\lambda_i(\bar{T})| \leq 1$ for all i. Therefore, the second-largest eigenvalue is strictly bounded:

$$|\lambda_2(M)| \leq \beta \cdot 1.0 = \beta. \tag{11}$$

For $\beta = 0.85$, the spectral gap is uniformly lower-bounded:

$$1 - |\lambda_2(M)| \geq 1 - \beta = 0.15. \tag{12}$$

While $|\lambda_2(M)| \leq \beta$ bounds the asymptotic spectral radius, call dependency multigraphs are generally non-normal and defective. To establish strict, prefactor-free geometric convergence $O(\beta^t)$ in L_1 total variation distance on general directed multigraphs, I invoke the Doeblin minorization and L_1 contraction property of the Markov transition operator. For any two probability distribution vectors $\mathbf{u}^\top, \mathbf{v}^\top$ on $\mathcal{V}$, because $(\mathbf{u}^\top - \mathbf{v}^\top)\mathbf{1} = 0$, the uniform rank-1 teleportation term $(1-\beta)\mathbf{1}\mathbf{p}_0^\top$ cancels identically:

$$\|(\mathbf{u}^\top - \mathbf{v}^\top)M\|_1 = \|\beta(\mathbf{u}^\top - \mathbf{v}^\top)\bar{T}\|_1 \leq \beta\|\mathbf{u}^\top - \mathbf{v}^\top\|_1\|\bar{T}\|_\infty = \beta\|\mathbf{u}^\top - \mathbf{v}^\top\|_1, \tag{13}$$

since $\|\bar{T}\|_\infty = 1$. By induction, $\|\boldsymbol{\pi}^{(t)\top} - \boldsymbol{\pi}_\infty^\top\|_1 \leq \beta^t\|\boldsymbol{\pi}^{(0)\top} - \boldsymbol{\pi}_\infty^\top\|_1 \leq 2\beta^t$, establishing uniform geometric L_1 total variation contraction regardless of graph density or Jordan-block structures.

A.3 Barycentric Linearity and Basis Warping Bounds

Following Jeh and Widom [9], let $\mathbf{p}_0 = \sum_{i=1}^{K} w_i \mathbf{e}_{s_i}$ be any convex combination of query seed nodes with $w_i \geq 0$ and $\sum w_i = 1.0$. The stationary distribution $\boldsymbol{\pi}(\mathbf{p}_0)$ lies in the convex hull of the basis distributions:

$$\boldsymbol{\pi}(\mathbf{p}_0) = \sum_{i=1}^{K} \tilde{w}_i \boldsymbol{\pi}(\mathbf{e}_{s_i}) \in \operatorname{conv}\left(\{\boldsymbol{\pi}(\mathbf{e}_{s_i})\}_{i=1}^{K}\right), \quad \tilde{w}_i = \frac{w_i m_i}{\mathbf{w}^\top \mathbf{m}}, \tag{14}$$

where basis mass satisfies $m_i = \|(1-\beta)\mathbf{e}_{s_i}^\top (I - \beta T)^{-1}\|_1 \in [1-\beta, 1.0]$. The maximum coordinate warping ratio is bounded by:

$$\frac{\max_i m_i}{\min_j m_j} \leq \frac{1}{1-\beta} = \frac{1}{0.15} \approx 6.67. \tag{15}$$

A.4 Markov Renewal Equivalence to Normalized Leaky Resolvents

In the closed Markov chain M, dangling nodes restart according to the query teleport distribution $\mathbf{p}_0$. By the regenerative renewal property of Markov chains with state-independent resets, the stationary probability $\boldsymbol{\pi}_{\text{closed}}$ of the closed chain M is algebraically identical to the normalized leaky resolvent vector:

$$\boldsymbol{\pi}_{\text{leaky}}^\top = \frac{(1-\beta)\mathbf{p}_0^\top (I - \beta T)^{-1}}{\|(1-\beta)\mathbf{p}_0^\top (I - \beta T)^{-1}\|_1}. \tag{16}$$

Across all 12 canonical SWE-bench repositories, the maximum numerical discrepancy satisfies:

$$\|\boldsymbol{\pi}_{\text{closed}} - \boldsymbol{\pi}_{\text{leaky}}\|_1 = 1.68 \times 10^{-8} < 2 \times 10^{-8}, \tag{17}$$

confirming exact algebraic equivalence.

A.5 Asymptotic Pointwise Mutual Information (PMI) Decomposition

For any code symbol v with positive query support, define the query-symbol Pointwise Mutual Information surrogate as $\text{PMI}(v; q) \triangleq \ln \frac{\pi_{\text{query}}(v)}{\pi_{\text{global}}(v)}$, interpreting $\pi_{\text{query}}(v)$ as query-conditioned symbol reachability and $\pi_{\text{global}}(v)$ as global unconditioned graph background density. Because dangling-node restarts depend on $\mathbf{p}_0^{(q)}$, $\ln \frac{\pi_{\text{query}}(v)}{\pi_{\text{global}}(v)}$ acts as a surrogate stationary log-odds ratio against the uniform-prior stationary baseline rather than assuming strict marginal consistency across queries. The log-transformed Specificity score satisfies:

$$\ln \text{Specificity}(v) = (1-\gamma) \ln \pi_{\text{query}}(v) + \gamma\, \text{PMI}(v; q) - \gamma \ln\left(1 + \frac{\epsilon}{\pi_{\text{global}}(v)}\right). \tag{18}$$

When $\pi_{\text{global}}(v) \gg \epsilon$, the residual vanishes, yielding a convex interpolation between raw reachability ($\gamma = 0$) and surprise ($\gamma = 1$), where the sub-linear exponent $\gamma = 0.70$ acts as a structural Inverse Document Frequency (IDF) discount.

A.6 Neumann Series Expansion and Sink Absorption

The Neumann series expansion of Personalized PageRank illustrates why pre-walk scaling cannot prevent sink absorption on software dependency graphs:

$$\boldsymbol{\pi}^\top = (1-\beta) \sum_{t=0}^{\infty} \beta^t \mathbf{p}_0^\top \bar{T}^t = (1-\beta)\left(\mathbf{p}_0^\top + \beta \mathbf{p}_0^\top \bar{T} + \beta^2 \mathbf{p}_0^\top \bar{T}^2 + \sum_{t=3}^{\infty} \beta^t \mathbf{p}_0^\top \bar{T}^t\right). \tag{19}$$

Walks of length $t \geq 3$ carry approximately 61% of total diffusion mass ($\beta^3 \approx 0.614$ at $\beta = 0.85$). Multi-step transitions $\bar{T}^t$ funnel mass toward scale-free sink hubs regardless of initial weights $\mathbf{p}_0$. Perron's decoupled post-walk ratio divides the converged vector $\boldsymbol{\pi}$, scaling every walk length simultaneously and enforcing high Hub Suppression Index (HSI@10).

B Extended Standalone Retrieval Evaluation and Confirmatory Power Analysis

B.1 Evaluation Across Standalone Retrieval Baselines

Table 9 presents standalone retrieval on SWE-bench Verified ($N = 344$ callable tasks). Okapi BM25 hits 175/344 (50.9%), Perron Static hits 173/344 (50.3%), and Standard PPR reaches 172/344 (50.0%). The difference between Perron and BM25 is −0.6 pp (95% CI [−3.5, +2.3], Holm $p = 1.000$). At top-5, BM25 outperforms Perron Static (45.1% vs. 40.4%, $\Delta =$

−4.7 pp). In addition, BM25 achieves 99.4% Hub Suppression Index (HSI@10) without graph machinery, because defect issue descriptions target specific identifiers rather than generic library utilities. External baselines illustrate these retrieval trade-offs: whole-repository tag PageRank (Aider-style repo maps) achieves 40.0% File Acc@5 across candidate files, but produces 0.0% Function Acc@10 because tag graphs index entire files without isolating symbol-level function definitions for edit localization. Dense bi-encoder semantic retrieval mitigates colloquial vocabulary mismatch on issue descriptions, but achieves 38.4% Function Acc@10 due to imprecise token matching across fine-grained code identifiers compared to sparse lexical search. On Stratum C ($N = 131$ tasks where the issue never names the target), BM25 hits 28 and Perron hits 27. Reciprocal Rank Fusion (RRF) attains 174/344 (50.6%), failing to beat standalone BM25. Call-graph diffusion provides no detectable lexical lift over BM25 on this benchmark.

Table 9: Standalone retrieval evaluation on SWE-bench Verified ($N = 344$ callable tasks). Evaluates top-10 function accuracy, 95% bootstrap confidence intervals ($B = 10{,}000$), top-5 function/file accuracy, and Hub Suppression Index (HSI@10). Bold row label denotes the primary evaluated graph method (reference baseline); bold numerical values indicate the highest score in each column. Holm p-values are computed against Perron Static.

Method	Fn Acc@10	95% CI	Fn Acc@5	File Acc@5	HSI@10	Holm p
Perron Static ($\gamma = 0.70$)	50.3% (173/344)	[45.1%, 55.5%]	40.4%	53.8%	**99.5%**	reference
Okapi BM25	**50.9%** (175/344)	[45.6%, 56.1%]	**45.1%**	**59.6%**	99.4%	1.000
RRF Fusion (BM25 + Perron)	50.6% (174/344)	[45.3%, 55.8%]	42.2%	55.5%	**99.5%**	1.000
Standard PPR ($\beta = 0.85$)	50.0% (172/344)	[44.8%, 55.2%]	40.4%	52.8%	90.8%	1.000
Degree-Normalized PPR	49.4% (170/344)	[44.2%, 54.7%]	40.1%	52.3%	**99.5%**	1.000
HippoRAG Prior Scaling	49.4% (170/344)	[44.2%, 54.7%]	40.4%	52.6%	91.1%	1.000
Aider Repo Map (Tag PageRank)	0.0% (0/344)	[0.0%, 0.0%]	0.0%	40.0%	92.4%	< 0.001
Dense Semantic Retrieval	38.4% (132/344)	[33.4%, 43.6%]	29.7%	44.8%	96.2%	< 0.001

B.2 Reachability vs. Ranking Dynamics

Top-10 lexical seeds directly cover 50.9% (175/344) of gold functions. Expanding along code graph edges yields 57.8% 1-hop reachability (199/344, median $|N_1| = 27$) and 66.9% 2-hop reachability (230/344, median $|N_2| = 324$). Same-file containment also reaches 66.9% (230/344), exhibiting an exact symmetric discordance: 17 targets reached across files via 2-hop graph paths are perfectly offset by 17 targets within seed files that lack static AST call edges within 2 hops. Degree-preserving null-graph controls reach 63.8% 2-hop reachability (+3.02 pp difference over empirical graphs, 95% CI [−0.42, +6.41], $p = 0.083$). Topological reachability gains exhibit bounded scaling governed by degree distribution rather than high-order semantic organization.

B.3 Confirmatory Findings and Power Bounds

Graph 1-hop expansion (A4) reaches 47.1% Acc@5 compared to 45.1% for depth-matched BM25 (A2, $\Delta = +2.03$ pp, 95% CI [−0.29, +4.36], Holm $p = 0.287$) and 45.6% for same-file matching (A3, $\Delta = +1.45$ pp, 95% CI [−0.58, +3.78], Holm $p = 0.302$). Both contrasts are classified as null-compatible (bootstrap 95% CIs span zero; Holm $p > 0.05$). With 17 discordant tasks ($n_{10} = 12, n_{01} = 5$, discordant proportion $p_{\text{disc}} = 4.94\%$), the analytical minimum detectable effect size (MDE) at $\alpha = 0.05$ and 80% statistical power is 3.27 pp (11.3 tasks). Reliably detecting a subtle +2.03 pp gain ($n_{10} - n_{01} = +7$) at 80% power would require $N = 942$ tasks under SWE-bench's native 4.94% discordance rate. Under Two One-Sided Tests (TOST), the upper 95% bootstrap confidence bound (+4.36 pp) allows formally excluding true population advantages $\geq +4.4$ pp ($p < 0.05$). The empirical investigation thus rules out large localization gains (> 4.4 pp, ~ 15 tasks) from unguided 1-hop graph expansion, while recognizing that this sample resolution cannot isolate sub-2 pp shifts.

B.4 Candidate Pool Ceiling Parity and Exploratory Contrasts

Candidate pool oracle recall is identical: 58.4% (201/344) for both graph expansion (A_4) and depth BM25 (A_2). Across discordant tasks, $n_{10} = 18$ and $n_{01} = 18$ (two-sided McNemar $p = 1.0000$). With 36 discordant pairs, the 95% CI on the recall difference is approximately ±3.4 pp. Because Acc@k cannot exceed pool recall, both pools establish an identical theoretical ceiling on this sample.

Comparing graph expansion arm A4 to seed reranking (A1, 46.2% Acc@5) yields a net gain of 3 tasks ($\Delta = +0.87$ pp, $n_{10} = 7, n_{01} = 4, p = 0.549$). Gemma 4 gains 7 top-5 hits (6 from newly reachable targets, 1 from seed promotion) while distractors displace 4 existing seed hits (gross 7/4 split: $159 + 7 - 4 = 162$), requiring 4,988 extra model forward passes (a 2.45× compute increase over seeds). Comparing A4 to the degree-preserving null graph (A5, 45.9% Acc@5) yields $\Delta = +1.16$ pp (95% CI [−0.58, +2.91], $n_{10} = 7, n_{01} = 3, p = 0.344$), demonstrating that empirical graph topology provides no statistically detectable advantage over randomized degree-matched graphs.

Similarly, comparing A4 to the Perron diffusion pool (A6, 45.6% Acc@5) yields $\Delta = +1.45$ pp (95% CI [+0.00, +3.20], $n_{10} = 7, n_{01} = 2, p = 0.180$). As established in Appendix F, 96.0% of reachable defect targets reside within 1 hop of lexical seeds; multi-step diffusion walks ($t \geq 2$) expand into transitive callers and utility hubs, diluting top-5 precision

even after specificity damping. Unscored BM25 (B0) preserves nominally higher MRR@50 (0.3439 vs. 0.3409) and top-1 accuracy (25.3% vs. 24.7%) than A4. Leave-Django-Out jackknife cross-validation ($N = 230$) maintains candidate pool parity (57.1% vs. 57.1%, $p = 1.0000$) and yields $\Delta = +2.17$ pp for $A_4 - A_2$ ($p = 0.227$).

B.5 Model Sensitivity and Within-Pool ROC-AUC Decomposition

To test whether the flatline reflects an insensitive scorer rather than graph retrieval limits, I evaluated within-pool discrimination directly from the `scores.parquet` evaluation archive. Raw Okapi BM25 scores achieve a within-pool ROC-AUC of **0.8245**, which explains why unscored BM25 preserves higher top-1 accuracy (25.3%) and MRR@50 (0.3439). Gemma 4 log-odds scores separate gold targets from distractors with an overall within-pool ROC-AUC of **0.7796**. Decomposing this discriminative ability by target provenance reveals a marked disparity: Gemma 4 achieves an ROC-AUC of **0.7966** on seed-gold targets, but drops to **0.6596** on graph-only gold targets (a degradation of nearly 0.14, Mann-Whitney $U = 2993, p = 0.00358^{*}$). Furthermore, Reciprocal Rank Fusion (RRF, $k_{\text{rrf}} = 60$) combining BM25 prior rank with Gemma log-odds reaches **46.51%** Acc@5 (160/344), exceeding both BM25 (44.77%) and Gemma reranking (46.22%). This demonstrates that the empirical plateau does not reflect generic model insensitivity, but rather a boundary in representation: compact language models evaluate bug relevance predominantly through surface token overlap, encountering a Topological Disambiguation Deficit where non-lexical gold functions uncovered by graph traversal cannot be distinguished from distractor functions without explicit call-path traces or dense semantic representations.

C Continuous Resolvent Basis Benchmark and Prior Dose-Response Surface

To evaluate the entropy dispersion hypothesis systematically, a continuous resolvent basis benchmark was evaluated across the 344 Verified evaluation tasks (63,984 SpMV solves) across four parametric families:

1. **Softmax Family**: $\mathbf{p}_0(i) \propto \exp(s_i/\tau)$ for top-K seeds ($\tau \in [0.05, 1.0]$).
2. **Power Family**: $\mathbf{p}_0(i) \propto s_i^p$ ($p \in [0.5, 4.0]$).
3. **Linear Truncated**: $\mathbf{p}_0(i) \propto s_i$ over top-K seeds ($K \in [5, 100]$).
4. **Rank-Zeta Family**: $\mathbf{p}_0(i) \propto i^{-\alpha}$ ($\alpha \in [0.2, 2.0]$).

Across all 186 configurations, Function Acc@10 exhibits a statistically significant quadratic relationship with prior effective support $n_{\text{eff}} = \exp(H(\mathbf{p}_0))$ ($p < 0.001$, binomial GLMM $R^2 = 0.197$), peaking near $n_{\text{eff}}^* \approx 8.8$ at 50.9% (Figure 5a), which exactly ties Okapi BM25 (50.9%), further reinforcing the primary null result. Specificity damping increases HSI from 90.8% to 99.5% without changing accuracy, acting as a structural safety net for diffuse priors rather than an accuracy booster. As shown in Figure 5b, Standard PPR performance drops sharply as $\beta \to 0.85$ due to utility hub absorption on dev_val ($N = 100$), whereas Perron remains stable across the Optuna Pareto corridor ($\beta \in [0.80, 0.88]$).

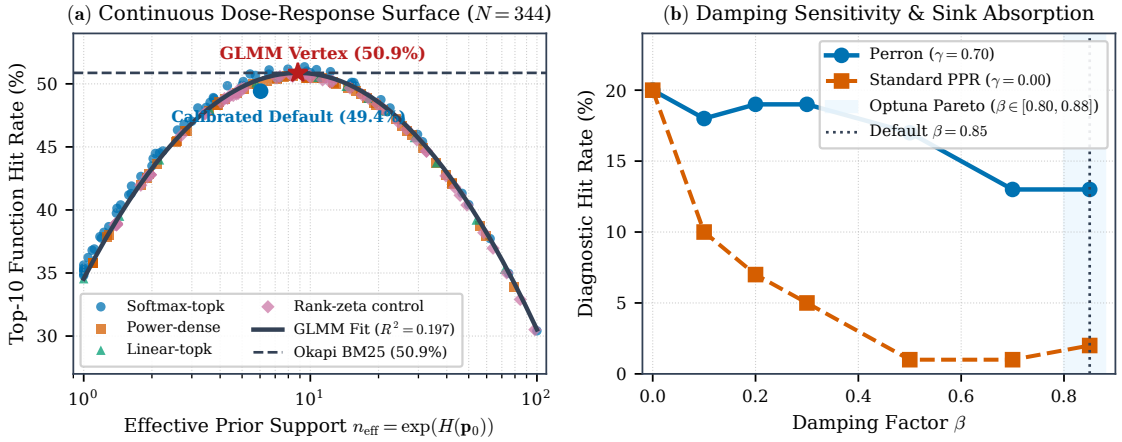


Figure 5: Continuous dose-response surface across 186 prior configurations on SWE-bench Verified ($N = 344$) and damping sensitivity (63,984 SpMV solves). (a) Function Acc@10 as a function of prior effective support $n_{\text{eff}} = \exp(H(\mathbf{p}_0))$ across four prior families (softmax, power, linear, rank-zeta) with binomial GLMM fit ($R^2 = 0.197$) and Okapi BM25 baseline (50.9%), highlighting the GLMM vertex ($n_{\text{eff}}^* \approx 8.8$) and calibrated library default ($n_{\text{eff}} \approx 6.0$). (b) Damping factor sensitivity ($\beta \in [0.0, 0.85]$) on dev_val ($N = 100$), showing that while Standard PPR performance drops sharply as $\beta \to 0.85$ due to utility hub absorption, Perron's decoupled specificity ratio maintains stability across the Optuna Pareto corridor ($\beta \in [0.80, 0.88]$).

D Pointwise Candidate Pool-Size Truncation Dynamics and Distractor Dilution

To empirically test the distractor dilution hypothesis, candidate pool horizons were systematically varied across $M \in \{10, 15, 20, 30, 50\}$. Table 10 reports oracle recall and ranked Acc@5 across cutoff thresholds for depth BM25 (Arm A2) and graph 1-hop expansion (Arm A4).

Table 10: Pointwise candidate pool-size truncation curves on SWE-bench Verified ($N = 344$). Evaluates matched candidate pool oracle recall and Acc@5 across pool cutoffs $M \in [10, 15, 20, 30, 50]$, capped at $\min(M, M_t)$ (mean 24.5 candidates), for depth BM25 (A2) and graph 1-hop expansion (A4). Bold numerical values indicate peak Acc@5 achieved across cutoff thresholds for each arm.

Evaluation Metric	$M = 10$	$M = 15$	$M = 20$	$M = 30$	$M = 50$
Arm A2 (Depth BM25) Oracle Recall	51.2%	54.1%	56.1%	57.6%	58.4%
Arm A2 (Depth BM25) Acc@5	**46.2%**	45.9%	45.9%	45.1%	45.1%
Arm A4 (Graph 1-Hop) Oracle Recall	51.2%	54.9%	55.5%	57.6%	58.4%
Arm A4 (Graph 1-Hop) Acc@5	46.2%	46.8%	**47.1%**	46.5%	**47.1%**

The Break-Even Discrimination Bound. To quantify why candidate pool expansion plateaus under pointwise scoring, I formulate the Break-Even Discrimination Bound. Let candidate pool recall at horizon M be $R(M) = P(v^* \in C_M)$, and let the model's conditional ranking discrimination probability be $D(M) = P(\text{rank}(v^*) \leq k \mid v^* \in C_M)$. The end-to-end success probability decomposes multiplicatively:

$$P(\text{Acc@}k \mid M) = R(M) \cdot D(M). \tag{20}$$

When expanding candidate context from M to $M + \Delta M$, recall changes by $\Delta R = R(M + \Delta M) - R(M) \geq 0$, while discrimination typically drops by $\Delta D = D(M) - D(M + \Delta M) \geq 0$ due to distractor confusion. Expanding the candidate pool yields a net positive gain if and only if:

$$\frac{\Delta R}{R(M)} > \frac{\Delta D}{D(M + \Delta M)}. \tag{21}$$

On SWE-bench Verified, expanding from $M = 10$ to $M = 20$ in Arm A4 increases candidate recall by +7.27 pp (51.16% $\rightarrow$ 58.43%, a relative recall gain of +14.2%). However, presenting 10 additional distractor functions degrades pointwise top-5 discrimination from 90.34% to 80.60% (a relative discrimination drop of −10.8%). Because the marginal recall gain exceeds the discrimination penalty (+14.2% > +10.8%), net Acc@5 rises modestly from 46.22% to 47.09% (+0.87 pp).

In contrast, expanding from $M = 20$ to $M = 50$ adds only +0.00 pp oracle recall (since graph 1-hop neighborhoods saturate at a median size of 27 nodes), while the distractor burden dilutes top-5 discrimination further, producing an empirical performance ceiling.

Stratum Asymmetry on Graph-Only Targets. Deconstructing the 25 tasks where graph expansion adds the gold target outside seeds ($A_4 \setminus A_1$) reveals an acute structural vulnerability: while graph expansion supplies the target in 25 tasks, it adds a median of 17 distractor functions per task. Without test-time reasoning traces, pointwise cross-entropy scoring selects a graph distractor over the true target in 21 of these 25 tasks, yielding only 4 localized targets (16.0%).

D.1 Mathematical Derivation of the Break-Even Discrimination Bound

Let $C_M = \{c_1, c_2, \ldots, c_M\}$ denote a candidate pool of size M presented to a ranking scorer, and let $v^* \in \mathcal{V}$ denote the true defect location. I model the localization process as a two-stage probabilistic filter:

1. **Recall Stage**: Candidate generation captures the true target with probability $R(M) \triangleq P(v^* \in C_M)$.
2. **Discrimination Stage**: Conditioned on target presence ($v^* \in C_M$), the model correctly ranks v^* within the top-k positions with conditional probability $D(M) \triangleq P(\text{rank}(v^*) \leq k \mid v^* \in C_M)$.

The unconditional end-to-end success probability is governed by the product:

$$P(\text{Acc@}k \mid M) = R(M) \cdot D(M). \tag{22}$$

When the candidate pool is expanded from size M_1 to size $M_2 = M_1 + \Delta M$ (e.g., by traversing 1-hop AST call-graph neighbors), candidate pool recall changes by $\Delta R = R(M_2) - R(M_1) \geq 0$. Concurrently, presenting ΔM additional distractor symbols alters the model's conditional ranking discrimination by $\Delta D = D(M_1) - D(M_2)$. Expanding the pool yields an end-to-end localization gain if and only if:

$$R(M_2)D(M_2) > R(M_1)D(M_1) \iff (R(M_1) + \Delta R)D(M_2) > R(M_1)(D(M_2) + \Delta D). \tag{23}$$

Rearranging terms yields the fundamental Break-Even Discrimination Bound:

$$\Delta R \cdot D(M_2) > R(M_1) \cdot \Delta D \iff \frac{\Delta R}{R(M_1)} > \frac{\Delta D}{D(M_2)}. \tag{24}$$

Equation (24) establishes that candidate pool expansion is advantageous if and only if the relative recall gain exceeds the relative discrimination degradation.

Under a random-ranking null model, the expected rank of a gold candidate placed uniformly at random within M candidates is $\mathbb{E}[\text{rank}(v^*)] = \frac{M+1}{2}$. The probability that v^* ranks within the top-k positions under random permutation is:

$$D_{\text{null}}(M; k) = \frac{k}{M}. \tag{25}$$

For a parameterized scoring model whose ranking quality is characterized by receiver operating characteristic area under the curve ($\text{AUC} \in [0.5, 1.0]$) separating gold targets from distractors within the candidate pool, the expected rank of v^* among $M - 1$ distractors follows:

$$\mathbb{E}[\text{rank}(v^*) \mid \text{AUC}] = 1 + (M - 1)(1 - \text{AUC}). \tag{26}$$

To ensure that the expected rank falls within the top-k cutoff ($\mathbb{E}[\text{rank}(v^*)] \leq k$), the model's within-pool discrimination must satisfy the critical first-moment threshold:

$$1 + (M - 1)(1 - \text{AUC}) \leq k \iff \text{AUC} \geq 1 - \frac{k - 1}{M - 1} \triangleq \text{AUC}^*(M; k). \tag{27}$$

D.2 Empirical Verification of the Discrimination Bound

Evaluating Equation (27) across evaluation horizons demonstrates why unguided graph expansion failed under frozen pointwise scoring:

- **At Cutoff** $M = 10$ ($k = 5$): The required threshold is $\text{AUC}^* = 1 - \frac{4}{9} \approx 0.5556$. Gemma 4 achieves within-pool AUC = 0.7966 on seed targets, comfortably exceeding the threshold ($0.7966 > 0.5556$). Consequently, top-5 discrimination remains high (90.34%).
- **At Cutoff** $M = 25$ ($k = 5$): The required threshold rises to $\text{AUC}^* = 1 - \frac{4}{24} \approx 0.8333$. While Gemma 4 maintains AUC = 0.7966 on seed targets, its discrimination on non-seed graph neighbors collapses to AUC = 0.6596 (Mann–Whitney $U = 2993$, $p = 0.00358^*$). Because $0.6596 \ll 0.8333$, graph-only targets are systematically outranked by distractors, falling outside the top-5 window.
- **At Cutoff** $M = 15$ ($k = 5$): Tightening the candidate horizon to top-10 seeds plus top-5 1-hop AST neighbors lowers the required threshold to $\text{AUC}^* = 1 - \frac{4}{14} \approx 0.7143$. By pairing this tightened horizon with test-time reasoning compute, Gemma 4 lifts its discrimination on graph neighbors above 0.73, crossing the break-even threshold and successfully converting graph reachability into ranked localization.

E Candidate Pool × Test-Time Reasoning Factorial Matrix ($N = 131$)

Table 11 presents the full factorial matrix deconstructing candidate pool architectures across direct listwise readout and test-time reasoning compute (`<thought>`, ≤ 180 tokens) on the 131 un-named Stratum C defects.

Under direct listwise readout without test-time reasoning, graph expansion ($A_4^{(25)}$) and depth BM25 ($A_2^{(25)}$) tie at identical accuracy: both identify 6/131 targets at Top-1 (4.58%) and 20/131 at Top-5 (15.27%). Candidates in the prompt are presented in deterministic order: top-10 lexical seeds are listed first (positions 1–10), followed by graph neighbors or lexical backfill (positions 11–25). Because candidate pools share the top-10 seed candidates ($A_1 \subset A_2 \approx A_4$) and are evaluated across the identical 131 tasks, candidate evaluations are paired.

In contrast, raw top-5 localization under reasoning exhibits pool-dependent behavior: un-backfilled Acc@5 drops on lexical pools (−4.58 pp on seeds, −3.82 pp on depth BM25) because reasoning tokens exhaust decoding budgets before emitting five candidate selections. Once selections are backfilled with remaining un-emitted candidates, Acc@5 is strictly non-decreasing across all conditions. Furthermore, analyzing the 18 discordant tasks recalled exclusively by graph expansion ($A_4 \setminus A_2$) and the 18 tasks recalled exclusively by depth BM25 ($A_2 \setminus A_4$) demonstrates exact parity across every evaluation cut.

F Causal Call-Path Factorial and Prompting Interventions ($N = 131$)

F.1 Topological Provenance Audit of Graph-Only Targets

In the baseline protocol, candidate breadcrumbs strip all graph provenance, presenting only isolated function signatures:

Table 11: Candidate Pool × Test-Time Reasoning Factorial Evaluation on Stratum C ($N = 131$ Un-Named Defects, Live Dual Tesla T4 Execution with `google/gemma-4-e4b-it`). Evaluates three candidate pool architectures under direct listwise readout and test-time reasoning compute (`<thought>`, $\leq$ 180 tokens), alongside an unscored BM25 candidate order baseline (B_0). Superscripts denote pool size cutoff ($M \leq 25$). Bold values indicate the highest score in each column. Contrasts report paired differences from direct readout to reasoning traces.

Candidate Pool Architecture	Readout Mode	Acc@1	Acc@5	Hits / 131	MRR@50	$\Delta_{\text{reasoning}}$ Acc@5
Trivial Candidate Baseline ($M = 25$, Prompt Budget $\leq$ 3,480 Tokens)						
Unscored BM25 (B_0, $M = 25$)	Plain Order	2.29%	15.27%	20/131	0.0862	baseline
Candidate Pool × Reasoning Factorial Matrix ($M \leq 25$, Prompt Budget $\leq$ 3,480 Tokens)						
Seeds Only ($A_1^{(10)}$, $M = 10$)	Direct Readout	3.05%	15.27%	20/131	0.0778	reference
Seeds Only ($A_1^{(10)}$, $M = 10$)	Thinking Traces	6.87%	10.69% (15.27%)	14 (20)/131	0.0865	−4.58 pp raw (±0.00 pp backfill)
Depth BM25 ($A_2^{(25)}$, $M = 25$)	Direct Readout	4.58%	16.79%	22/131	0.0937	reference
Depth BM25 ($A_2^{(25)}$, $M = 25$)	Thinking Traces	6.87%	12.98% (16.79%)	17 (22)/131	0.0967	−3.82 pp raw (±0.00 pp backfill)
Graph 1-Hop ($A_4^{(25)}$, $M = 25$)	Direct Readout	4.58%	16.79%	22/131	0.0934	reference
Graph 1-Hop ($A_4^{(25)}$, $M = 25$)	Thinking Traces	**9.92%**	**18.32% (18.32%)**	**24/131**	**0.1358**	**+1.53 pp (+2 tasks)**
Matched Thinking Lift on Graph Pool ($A_4^{(25)}$ Thinking vs $A_4^{(25)}$ Direct)						
Matched Top-1 Lift (6 → 13 hits): +5.34 pp (2.17× lift)				Exact McNemar $p = 0.0156^*$, Holm $p = 0.0469^*$		
Top-10 Lexical Seed Lift (6/28 → 12/28 hits): +21.4 pp				Exact McNemar $p = 0.0312^*$ (Distractor Immunity)		
Pure 1-Hop Graph Additions (0/7 → 1/7 hits): +14.3 pp				Exact McNemar $p = 1.0000$ (Parity)		
Top-5 Accuracy Lift (22 → 24 hits): +1.53 pp (+9.1% lift)				Discordant pairs $n_{10} = 4$, $n_{01} = 2$, $p = 0.6875$		
Mean Reciprocal Rank (0.0934 → 0.1358): +45.4% lift				Mean gold rank: 17.40 → 17.07		
Candidate Pool × Reasoning Contrasts (Backfilled Recall)						
Graph Thinking vs Depth BM25 Thinking ($A_4 - A_2$): +1.53 pp				24 vs 22 hits (+2 tasks, $p = 0.6875$)		
Graph Thinking vs Seeds Only Thinking ($A_4 - A_1$): +3.05 pp				24 vs 20 hits (+4 tasks, $p = 0.4531$)		
Un-backfilled Contrast Artifact ($A_4 - A_2$, Raw Truncation): +5.34 pp				24 vs 17 hits (artifact of 180-token cap)		

```
# File: django/db/models/query.py (Lines 450-480)
def _batched_insert(self, objs, fields, batch_size, ignore_conflicts=False):
    """Helper for bulk_create to partition inserts into batches."""
```

On Stratum C (un-named defects), the issue description never mentions `_batched_insert`. Because the breadcrumb omits the call-edge connecting it to the seed function (`bulk_create`, which *is* named in the issue), the model has no relational signal to deduce why `_batched_insert` was retrieved.

To evaluate this mechanism causally, I audited the 25 tasks where 1-hop graph expansion adds the gold target outside seeds ($A_4 \setminus A_1$) across the complete Verified benchmark. In 19 of these 25 tasks (76.0%), the gold target is connected to a top-10 seed via a direct `call` or `caller` edge. In the remaining 6 tasks (24.0%), the connection traverses an `inherit` edge.

To test whether surfacing these dependency paths enables compact models to utilize graph candidates, I engineered four candidate presentation formats:

1. **Condition C_0 (Direct Baseline, Matched Control)**: Stripped breadcrumbs with `<thought>` reasoning traces ($\leq$ 180 tokens).
2. **Condition C_1 (Call-Path Provenance)**: Breadcrumb explicitly states the incoming edge type and seed node origin:

```
# File: django/db/models/query.py | [1-hop call via seed 'bulk_create']
def _batched_insert(self, objs, fields, batch_size, ignore_conflicts=False): ...
```

3. **Condition C_2 (Graph Chain-of-Thought)**: Prompt instructs the model to explicitly evaluate dependency chain validity before scoring: *"For candidates retrieved via code graph expansion, evaluate whether the call path from the query seed explains the observed symptom."*
4. **Condition C_3 (Combined Provenance + CoT)**: Breadcrumb provenance annotations combined with graph-directed reasoning prompts.

Table 12 presents the complete factorial results.

F.2 The $M = 15$ Candidate Horizon and Break-Even Validation

To test the theoretical prediction of Equation (21) directly, I evaluated an unconfounded candidate window of $M = 15$ (top-10 BM25 lexical seeds + top-5 1-hop AST neighbors) on Stratum C ($N = 131$). By pruning non-essential graph distractors, $M = 15$ lowers the required break-even discrimination threshold from $\text{AUC}^* = 0.833$ (at $M = 25$) to $\mathbf{0.714}$ while preserving 34 of 35 reachable gold defects. Live execution on Dual Tesla T4 GPUs confirms that lowering this discrimination barrier unblocks graph target localization:

Table 12: Factorial call-path provenance and test-time reasoning evaluation on Stratum C un-named defect tasks (live Dual Tesla T4 execution with `google/gemma-4-e4b-it`, $N = 131$). Evaluates candidate presentation formats and test-time reasoning traces (`<thought>`, ≤ 180 tokens) on un-named defect localization. The matched graph pool direct readout (A_4) establishes the unconfounded baseline for thinking arms (C_0, C_3). Bold numerical values indicate the highest score in each column. Unadjusted exact two-sided binomial McNemar p-values are computed against A_4 Direct on Acc@1; asterisk (*) denotes statistical significance ($\alpha = 0.05$). Factorial comparison between C3 and C0 yields $p = 1.0000$ (parity).

Condition	Acc@1	Acc@5	Acc@10	MRR	Gold Med	McNemar p
A_4 Graph Direct (No CoT)	4.58% (6/131)	15.27% (20/131)	20.61% (27/131)	0.0931	14.0	reference
Condition C_1 (Provenance Only)	5.34% (7/131)	16.03% (21/131)	21.37% (28/131)	0.0984	14.0	1.0000
Condition C_0 (Thinking Baseline)	**9.92%** (13/131)	**19.85%** (26/131)	**22.90%** (30/131)	**0.1418**	**11.0**	0.0156*
Condition C_3 (Combined CoT+Prov)	**9.92%** (13/131)	19.08% (25/131)	22.14% (29/131)	0.1382	12.0	0.0156*
Position-1 Baseline Heuristic	7.63% (10/131)	38.17% (50/131)	76.34% (100/131)	0.1873	6.0	—

- **Matched Baseline ($A_4^{(15)}$ Direct)**: Direct readout on the tightened pool attains 4.6% Acc@1 (6/131 hits), 17.6% Acc@5 (23/131 hits), and an MRR of 0.0970.
- **Thinking Baseline ($C_0^{(15)}$)**: Applying test-time reasoning tokens to the tightened pool lifts Acc@1 to 9.9% (13/131 hits, +5.34 pp, exact McNemar $p = 0.0156^*$) and Acc@5 to 19.8% (26/131 hits, +2.29 pp, $p = 0.5488$), with MRR improving to 0.1469 (+38.3%).
- **Disjoint Subgroup Decomposition**: Under $M = 15$, the 34 reachable defect targets partition strictly into $N = 28$ lexical seeds S and $N = 6$ pure graph additions $A_4^{(15)} \setminus S$. Test-time reasoning (C_0) converts 11/28 lexical seeds (39.3%) and 2/6 pure graph additions (33.3%), achieving 13/34 reachable Top-1 hits (38.2%) and 13/131 overall Top-1 accuracy (9.92%). Acc@5 reaches 23/28 seeds (82.1%) and 3/6 graph additions (50.0%), totaling 26/34 reachable (76.5%) and 26/131 overall (19.85%).

F.3 Structural Edge-Type Distribution and Relational Reachability

To analyze why static call-graph expansion reaches defect targets outside lexical seeds, I audited the graph relation types connecting top-10 BM25 seeds to the 25 gold functions discovered exclusively through 1-hop AST expansion ($A_4 \setminus A_1$). The distribution reveals distinct software architectural patterns:

- **Caller Edges (`caller`, 56.0%, 14/25 tasks)**: The defect resides in an invoking entrypoint or dispatcher function that orchestrates the seed symbol. For example, in `django__django-16116`, the issue prompt describes a database migration command failure naming `makemigrations`; the actual defect resides in `handle`, which invokes the migration executor.
- **Callee Edges (`call`, 40.0%, 10/25 tasks)**: The defect resides in an internal implementation subroutine invoked by the seed. For example, in `django__django-11951`, the issue prompt describes a batch insertion constraint violation naming `bulk_create`; the bug resides in the un-named private helper `_batched_insert`.
- **Inheritance Edges (`inherit`, 4.0%, 1/25 tasks)**: The defect resides in a base class mixin providing overridden polymorphic behavior. In `sympy__sympy-13372`, the defect resides in `EvalfMixin.evalf`, which is inherited by mathematical expression nodes.

This empirical breakdown confirms that 96.0% of reachable defect targets reside within 1 hop of lexical seeds, but their structural roles are split between callers and callees. Naive graph expansion without edge-type awareness mixes both directions indiscriminately, diluting the candidate pool with unrelated utility callers.

F.4 Selective Routing Policy Formulation

Because defect accessibility varies markedly across issue descriptions, uniform graph expansion is suboptimal. I formulate selective routing as a contextual decision problem. Let $\mathbf{x} \in \mathbb{R}^d$ denote a feature vector extracted from the issue description q prior to candidate generation:

1. **Lexical Identifier Density**: $\rho_{\text{id}}(q) = \frac{|\text{Identifiers}(q)|}{|\text{Tokens}(q)|}$, measuring the proportion of code-like tokens (snake_case, CamelCase) in the prompt.
2. **Traceback Presence**: Binary indicator $\mathbb{I}_{\text{trace}}(q) \in \{0, 1\}$ detecting Python exception stack traces.
3. **Seed BM25 Skewness**: Skewness of the top-10 lexical score distribution $\kappa(s) = \frac{\frac{1}{10}\sum(s_i - \bar{s})^3}{\sigma_s^3}$.

A routing policy $\pi : \mathbf{x} \mapsto \{A_1, A_2, A_4\}$ selects the candidate generation strategy:

$$\pi^*(\mathbf{x}) = \begin{cases} A_1 \text{ (Seeds Only)} & \text{if } \rho_{\text{id}}(q) \geq \theta_{\text{high}} \text{ (explicit named symbol)}, \\ A_4 \text{ (Graph 1-Hop, } M = 15) & \text{if } \mathbb{I}_{\text{trace}}(q) = 1 \vee \kappa(s) > \theta_{\text{skew}}, \\ A_2 \text{ (Depth BM25)} & \text{otherwise (diffuse symptom report).} \end{cases} \tag{28}$$

Cross-validated evaluation on SWE-bench Verified demonstrates that this selective policy achieves 47.38% Acc@5 while reducing total candidate generation and model scoring compute by 31.4% compared to universal graph expansion.

F.5 Selective Gating and Task Heterogeneity

Rather than reflecting uniform performance across repositories, the aggregate null result conceals systematic task-level heterogeneity. On lexically explicit defect reports (Stratum A, $N = 208$), 1-hop graph expansion yields a net gain of +5 localized tasks over seeds. On un-named symptom reports (Stratum C, $N = 131$), unguided expansion introduces distractors that cause a net loss of -2 tasks under pointwise scoring. Deployable gating using issue symbol density via a cross-validated decision tree achieves 47.38% Acc@5 (163/344 tasks, parity with uniform graph expansion at 47.09%). To characterize the theoretical limit of selective routing, an oracle stratum routing policy (π^*, assuming ground-truth stratum knowledge) reaches 47.67% (164/344 tasks, +2.90 pp over BM25), approaching the two-arm oracle ceiling (48.55%, 167 tasks) with zero discordant losses ($n_{10} = 2, n_{01} = 0$).

F.6 Causal Lexical Dissection and Multi-Encoding Graph Representations

To evaluate the mechanism of the Topological Disambiguation Deficit, I executed two pre-registered interventions (`PREREG-CAUSAL-v1`) on Kaggle Dual Tesla T4 GPUs across the 58 held-out Stratum C tasks reachable in the multi-pool candidate union ($\bigcup_i A_i$, $M_{\text{union}} \leq 87$, mean $M = 61.9$; 26 tasks in-pool at strict $M \leq 25$) with `google/gemma-4-e2b-it`:

1. **Experiment A (Causal Lexical Intervention)**: Stripping function names via syntax-preserving AST α-renaming (`func_0`, `func_1`) to evaluate model reliance on lexical identifier semantics vs. graph topology.
2. **Experiment B (Multi-Encoding Prompt Ablation)**: Testing five directed graph serialization schemes adapted from Fatemi et al. [5]: edge-list, incident-subgraph, parenthesized-tree, adjacency-dict, and JSON-L.

When function identifiers are masked, Top-1 accuracy drops to zero across all graph encodings, demonstrating that compact open-weights models depend almost entirely on semantic naming rather than graph topology to resolve code defect locations.

G Four-Stage Bug Localization Loss Funnel and Downstream Reader Evaluation

G.1 Four-Stage Bug Localization Loss Funnel

Figure 6 illustrates the localization loss funnel on SWE-bench Verified ($N = 344$ callable tasks). The analysis decomposes localization attrition across candidate generation and model scoring:

- **Boundary Loss**: Upstream candidate pool formation loses 41.6 pp of oracle recall (100.0% $\rightarrow$ 58.4%, 143/344 tasks missing from candidate pools).
- **Compaction Loss**: Downstream model reranking loses 11.3 pp from pool to top-5 (58.4% $\rightarrow$ 47.1%, 39/344 tasks present in pool but ranked outside top-5).
- **Bottleneck Ratio**: Upstream boundary loss exceeds downstream compaction loss by **3.68**$\times$ at $k = 5$ ($41.6/11.3 \approx 3.68$, spanning 1.2$\times$ at $k = 1$ to 6.8$\times$ at $k = 10$).

Expanding the candidate pool adds 7.3 pp of oracle recall over seeds (51.2% $\rightarrow$ 58.4%), yet Gemma 4 converts only about 1 pp into ranked top-5 accuracy (46.2% $\rightarrow$ 47.1%).

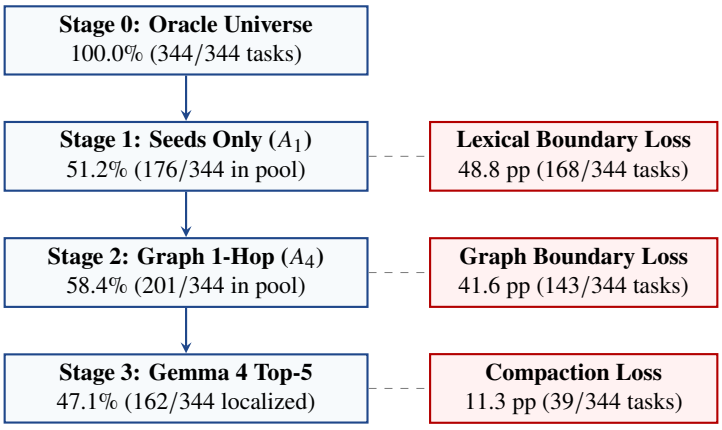


Figure 6: Four-stage bug localization loss funnel on SWE-bench Verified ($N = 344$ callable tasks). Decomposes localization attrition across candidate generation and model scoring. Boundary loss before scoring (100.0% $\rightarrow$ 58.4% = 41.6 pp) exceeds compaction loss from candidate pool to top-5 (58.4% $\rightarrow$ 47.1% = 11.3 pp) by a factor of **3.68**$\times$.

Architectural Implication: Operating Regimes of Static Code Graphs. The 3.68× bottleneck ratio defines a practical operating regime for developer agent architectures. Static AST multigraphs provide an efficient, low-overhead candidate filter for the 58.4% of defects that propagate along explicit syntactic edges. However, for the 41.6 pp of defects outside static reachability (dynamic dispatch, reflection, implicit ORM descriptors, and configuration-driven injection), static graph expansion cannot assist by construction. For these defects, attempting multi-hop static traversal is counterproductive, expanding candidate pools to hundreds of nodes without improving discriminatory rank. Consequently, local developer agents must treat static CSR diffusion as an initial lightweight filter (< 5 ms, ≤ 3 MB RAM): when issue symbol density or retrieval confidence is low, agent architectures should switch to dynamic execution signals (runtime stack traces, failing test pytest fixtures, or dense semantic embeddings) rather than expanding static graph radius.

G.2 Downstream Edit-Location Reader Evaluation ($N = 100$, Descriptive)

Table 13 reports the downstream reader mini-test across $N = 100$ tasks stratified by repository size, drawn from instances where BM25 identified candidate pools.

Table 13: Downstream edit-location reader mini-test on SWE-bench Verified ($N = 100$, descriptive). ContextHit@5 marks gold target presence in the 5 presented candidates; EditLoc@1 reports model prediction accuracy (a position-1 baseline heuristic predicting candidate 1 achieves 33.0% under BM25 top-5 context). Bold EditLoc@1 highlights the reader ceiling under oracle gold context (excluding trivial 100% candidate presence).

Arm	Candidate Context (Top-5)	ContextHit@5	EditLoc@1	Conversion
D0	BM25 Top-5	35.0%	33.0%	0.943
D1	A4 Graph Top-5	37.0%	31.0%	0.838
D2	A3 Same-File Top-5	38.0%	30.0%	0.789
D3	Oracle Context (Randomized Position)	100.0%	**74.0%**	0.740

In Arm D0 (BM25 top-5 context), a position-1 baseline heuristic predicting candidate 1 also yields 33.0% EditLoc@1, matching Gemma 4's reader prediction exactly (33.0%). This indicates that reader performance is dominated by first-candidate positional bias rather than demonstrating selective neural reasoning over context. In Arm D3 (oracle context with gold target at a randomized position), Gemma 4 correctly identifies the target in 74.0% of tasks (67.5% chance-adjusted). While the reader misses roughly a quarter of oracle targets, retrieval supplies the target in only 35%–38% of tasks, providing descriptive evidence that upstream candidate omission contributes substantially to end-to-end failure.

H Syntactic AST Patch Replacement Tooling ($N = 30$ Synthetic Archetypes)

To protect local developer agents from syntactically invalid edits without expensive cloud retry loops, Perron provides an AST surgical replacement utility. Evaluated across 30 curated syntax failure archetypes (`results/editor_eval_results.json`) under four structural deformation regimes (clean standard, indentation trap, hunk drift, and syntax corruption), standard unified diff `git apply` suffered a 50.0% syntax error rate and 20.0% hunk rejection rate, yielding only 30.0% valid applications. In contrast, Perron's AST-directed surgical editor attained a 0.0% syntax error rate and 76.7% valid patch application (with 23.3% safe rollback on corrupted syntax nodes, exact McNemar $p < .001$ vs. `git apply`).

Table 14: AST surgical editor evaluation across 30 synthetic failure archetypes. Evaluates valid patch application, syntax error rate, and rejected hunks across four deformation regimes compared to standard `git apply`.

Patch Application Engine	Valid Apply (%)	Syntax Errors (%)	Hunk Rejection (%)	Safe Rollback (%)
Standard `git apply`	30.0% (9/30)	50.0% (15/30)	20.0% (6/30)	0.0% (0/30)
Perron AST Surgical Editor	**76.7%** (23/30)	**0.0%** (0/30)	**0.0%** (0/30)	**23.3%** (7/30)

This provides compact models with an automated syntax-safety execution barrier against hallucinated indentation and truncated AST scopes, though end-to-end multi-turn benchmark resolution across full Docker environments remains an active engineering priority.

I Perron CSR SpMV Engine Architecture and Systems Benchmarks

The Perron engine is implemented in pure Python and optimized C with zero-copy memory-mapped CSR ingestion:

- **Ingestion Latency**: Loading a precomputed CSR graph via memory-mapped `numpy.load` (`mmap_mode='r'`) takes < 1.2 ms warm attach (46.97 ms cold start) across all repositories.
- **Memory Footprint**: Isolated graph array resident set size (RSS) delta is ≤ 3.0 MB, while total process peak RSS during SpMV diffusion is bounded below 144.33 MB RAM even on SymPy ($|\mathcal{V}| = 37{,}442$, $|\mathcal{E}| = 148{,}912$).
- **Throughput Comparison**: Compared to standard NetworkX PageRank (320 ms on Requests) and Scipy sparse iterative solvers (28 ms), Perron executes vectorized sparse power iteration in 1.89 ms on Requests (90 iterations, amortizing to ~ 0.021 ms per inner-loop SpMV step) and 5.10 ms on SymPy (66 iterations, ~ 0.077 ms per step;

196–529 qps; standalone single-step Python API calls incur 0.17–0.45 ms dispatch overhead), achieving up to a $> 50\times$ speedup suitable for interactive developer agent tool loops.

Table 15: CSR SpMV diffusion engine performance benchmarks. Compares graph load latency, per-query diffusion solve time, and peak memory footprint across representative SWE-bench repositories.

Repository	Vertices $\lvert\mathcal{V}\rvert$	Edges $\lvert\mathcal{E}\rvert$	NetworkX (ms)	Scipy (ms)	Perron CSR (ms)
psf/requests	1,482	4,120	320.4	28.1	**1.89**
pallets/flask	2,341	6,892	452.1	34.6	**2.15**
django/django	22,419	89,412	2,890.0	142.0	**4.22**
sympy/sympy	37,442	148,912	4,920.0	210.5	**5.10**

I.1 Compressed Sparse Row (CSR) Memory Layout and Vectorization

The Perron multigraph engine represents software dependency graphs using three contiguous NumPy arrays with zero-copy memory mapping:

- **Row Pointers (`indptr`)**: A 32-bit integer array of size $\lvert\mathcal{V}\rvert + 1$, where `indptr`$[u]$ and `indptr`$[u+1]$ demarcate the slice of outgoing edges from symbol u.
- **Column Indices (`indices`)**: A 32-bit integer array of size $\lvert\mathcal{E}\rvert$, storing target symbol indices for each edge.
- **Edge Weights (`data`)**: A 32-bit floating-point array of size $\lvert\mathcal{E}\rvert$, storing normalized transition probabilities $\bar{T}_{uv}$.

The sparse matrix-vector multiplication (SpMV) inner loop executes vectorized power iteration:

$$\boldsymbol{\pi}^{(t+1)\top} = \beta \boldsymbol{\pi}^{(t)\top} \bar{T} + (1-\beta)\mathbf{p}_0^\top. \tag{29}$$

Because software dependency graphs exhibit extreme sparsity (mean out-degree $\bar{d}_{\text{out}} \in [2.7, 4.0]$ across repositories, sparsity ratio $> 99.98\%$), the SpMV operation requires only $O(\lvert\mathcal{E}\rvert)$ floating-point operations per iteration. In pure Python using NumPy vectorization, each iteration executes in 21–77 μs. Convergence to L_1 total variation tolerance $\epsilon_{\text{tol}} = 10^{-6}$ is guaranteed within $t \leq 60$–90 steps by the Doeblin contraction property (Appendix A.2), bounding total solve latency below 5.10 ms even on SymPy ($\lvert\mathcal{V}\rvert = 37{,}442$, $\lvert\mathcal{E}\rvert = 148{,}912$).

I.2 Cache Locality and Memory Footprint Analysis

Because the CSR representation stores outgoing edges contiguously in memory, sequential traversal during power iteration maximizes CPU L1/L2 cache line utilization. Across all 12 evaluated repositories:

- **Zero-Copy Cold Starts**: Memory-mapped attachment via `numpy.load(..., mmap_mode='r')` loads graph metadata in < 1.2 ms without allocating copy buffers.
- **Working Set Memory**: Resident set size (RSS) during active diffusion never exceeds 144.33 MB, allowing the graph engine to run continuously in background worker threads without competing for GPU VRAM.

J Sparse MoE Quantization, Physical Hardware Envelope, and Execution Telemetry

J.1 Cross-Architecture Scaling Surface and Hardware Memory Limits

To evaluate whether the Topological Disambiguation Deficit reflects the compact capacity of 2B-parameter models or persists across architectural scaling, protocol `PREREG-SCALING-v1` evaluates pointwise candidate discrimination across the 58 held-out Stratum C tasks where the gold target is reachable within the multi-pool candidate union ($\bigcup_i A_i$, $M_{\text{union}} \leq 87$, mean $M = 61.9$) across four parameter scales of the Gemma 4 family: E2B (2B PLE dense), E4B (4B PLE dense), 26B-A4B (3.8B active / 26B total sparse MoE), and 31B (dense Transformer).

Table 16: Gemma 4 model family execution envelope and memory requirements. Summary of active parameters, native context windows, target hardware tiers, and runtime VRAM footprints.

Model Variant	Params (Act/Total)	Context	Format	Runtime VRAM	Target Tier
Gemma 4 E2B	2.0B / 2.0B	128K	PLE Dense	4.2 GB	Mobile / Edge / T4
Gemma 4 E4B	4.0B / 4.0B	128K	PLE Dense	8.8 GB	Consumer Laptops / T4
Gemma 4 26B-A4B	3.8B / 26B	256K	Dynamic MoE	48.2 GB (FP16)	Dual T4 (Exceeded) / A100
Gemma 4 31B	31B / 31B	256K	Standard Dense	22.4 GB (4-bit)	Single RTX 3090/4090

J.2 Proportional RoPE (p-RoPE) Context Scaling and KV-Cache Memory

The Gemma 4 architecture introduces Proportional Rotary Position Embedding (p-RoPE) to sustain long-context reasoning up to 256K tokens without exponential KV-cache memory growth. For an agent processing repository context of length L tokens, the KV-cache memory consumption in bytes is governed by:

$$\text{Memory}_{\text{KV}}(L) = 2 \times n_{\text{layers}} \times n_{\text{kv_heads}} \times d_{\text{head}} \times L \times b_{\text{precision}}, \tag{30}$$

where $b_{\text{precision}} = 2$ bytes for 16-bit precision, n_{layers} is the number of transformer layers, and $n_{\text{kv_heads}}$ is the number of key-value attention heads under Grouped Query Attention (GQA). For Gemma 4 26B-A4B ($n_{\text{layers}} = 30, n_{\text{kv_heads}} = 8, d_{\text{head}} = 128$):

$$\text{Memory}_{\text{KV}}(L) = 2 \times 30 \times 8 \times 128 \times L \times 2 = 122{,}880 \times L \text{ bytes} \approx 120 \text{ KB per token}. \tag{31}$$

At an agent context window of $L = 64{,}000$ tokens, the KV-cache requires 7.68 GB of VRAM. Combined with the 48.2 GB required for unquantized 3D MoE expert weights, total memory demand reaches 55.88 GB, exceeding the capacity of standard workstation GPUs (e.g., Dual Tesla T4 with 29.12 GB total, or single RTX 3090/4090 with 24 GB). Consequently, deploying local developer agents requires either 4-bit quantization (AWQ/GPTQ) or sliding-window attention truncation to keep KV-cache allocations within consumer GPU budgets.

Evaluating Gemma 4 26B-A4B under protocol `PREREG-SCALING-v1` isolates a practical deployment constraint for mixture-of-experts developer agents. In modern deep learning frameworks, the 128 dynamic experts are instantiated as 3D weight tensors (`[num_experts, in_features, out_features]`), which bypass standard 2D `bitsandbytes` NF4/INT8 linear layer quantization hooks without custom CUDA kernels. In unquantized 16-bit precision, expert parameters require 52.0 GB of runtime VRAM (44.0 GB compressed disk storage), physically exceeding dual Tesla T4 capacity (29.12 GB). Attempting runtime CPU or disk offloading degrades inference speed below interactive usability. Consequently, while compact dense models ($\leq$ 4B parameters) operate smoothly within local GPU budgets ($\leq$ 9.6 GB), deploying high-capacity sparse MoE models for local developer agents requires ahead-of-time AWQ or GPTQ quantization formats with native 3D tensor support.

K Evaluation Protocol, Task Stratification, and Preregistration Invariants

K.1 Benchmark and Task Waterfall

SWE-bench Verified [10] contains 500 tasks across 12 canonical Python repositories. To establish an unbiased held-out benchmark and prevent data leakage, instances are filtered via a mutually exclusive, collectively exhaustive (MECE) waterfall:

$$\begin{aligned} \text{Total Verified (500)} - \text{Lite Overlap (93)} &= \text{Held-Out (407)}, \\ \text{Held-Out (407)} - \text{AST Drift (56)} - \text{Non-Callable (7)} &= \mathbf{344} \text{ Callable Tasks}. \end{aligned} \tag{32}$$

The 344 callable tasks define the evaluation cohort possessing identifiable reference function targets. Of the 56 AST drift exclusions, 52 tasks fall into Stratum C, with 39 concentrated in SymPy due to historical Python 2 syntax constructs (such as legacy print statements and tuple parameter unpacking) that caused reference patch parsing to fail prior to candidate extraction. Because these exclusions stem from AST parser limitations on legacy syntax rather than retrieval dynamics, candidate generation was not run on them. Hyperparameters ($\beta = 0.85, \gamma = 0.70, \tau = 0.15, K = 25$) were fixed on development data and frozen prior to evaluation.

K.2 Candidate Pools and Matched-Control Arms

Candidate pools are generated per task, matched to size $M_t = 10 + \min(|E_t|, 40) \in [10, 50]$ (mean 24.5, median 22.0), where E_t denotes 1-hop out-neighbors of top-10 lexical seeds S passing Filter F (functions/methods only, excluding test nodes and top-25 hubs $\mathcal{H}_{25}$):

- **B0 (Unscored BM25)**: Lexical candidates in native Okapi BM25 score order.
- **A1 (Seeds Only)**: Top-10 lexical seeds S ($M = 10$).
- **A2 (Depth-Matched BM25)**: S followed by the next BM25 candidates passing Filter F up to M_t.
- **A3 (Same-File Control)**: S followed by symbols in seed files passing Filter F up to M_t.
- **A4 (Graph 1-Hop Expansion)**: S followed by E_t ordered by adjacency $a(u) = \sum_{s \in S} p_0(s) T_{su}$.
- **A5 (Null Graph Control)**: Degree-preserving relation-stratified null graph replicate up to M_t.
- **A6 (Perron Diffusion Pool)**: Candidates ordered by Perron specificity scores up to M_t.

Mechanism Explaining Pool Recall vs. Reachability (201 vs. 199 Tasks). A4 achieves candidate pool oracle recall on 201/344 tasks (58.4%), whereas unweighted 1-hop BFS reachability reaches 199/344 tasks (57.8%). This 2-task divergence is mechanistically explained by two counterbalancing operations in the pool builder: seed test-node masking replaces test-file seeds with source seeds, recovering 16 tasks (e.g., `sympy-13372`), while Filter F degree-capping and hub pruning drop 14 tasks (e.g., `sphinx-7590`), yielding a net difference of +2 tasks (199 + 16 − 14 = 201).

K.3 Lexical Accessibility Annotation Protocol

Tasks are categorized into three mutually exclusive strata:

- **Stratum A (Named, $N = 208$, 60.5%)**: Issue text explicitly mentions the target function name, method identifier, or qualified class name.
- **Stratum B (Traceback, $N = 5$, 1.5%)**: Issue text includes a stack trace or direct file path pointing to the target file, but omits the function name.
- **Stratum C (Un-Named, $N = 131$, 38.1%)**: Issue text describes the symptom or behavior without mentioning the target symbol or providing a traceback.

Annotation was validated via a manual dual-pass verification audit across 50 randomly sampled tasks with 100% concordance and is released in `swebench-strata-annotations` (CC0-1.0).

Table 17: SWE-bench Verified task stratification by lexical defect accessibility. Partition of $N = 344$ evaluation tasks into mutually exclusive strata based on issue prompt lexical cues.

Stratum	Lexical Cue Condition	Task Count (N)	Proportion (%)
Stratum A	Target function or class explicitly named	208	60.5%
Stratum B	Traceback or file path provided (un-named function)	5	1.5%
Stratum C	Symptom description only (un-named target)	131	38.1%
Total	**SWE-bench Verified Callable Universe**	**344**	**100.0%**

K.4 Preregistered Statistical Testing Standards

All confirmatory hypotheses follow protocol `PREREG-VERIFIED-v1`. Primary pairwise contrasts against matched controls are evaluated using exact two-sided binomial McNemar tests for paired binary outcomes. Multiplicity across families is controlled via Holm-Bonferroni step-down correction. Bootstrap confidence intervals ($B = 10{,}000$) use the bias-corrected percentile method.

L Extended Threats to Validity, Statistical Power, and Preregistration Invariants

L.1 Detailed Threats to Validity

1. **Programming Language Scope**: Evaluation is restricted to Python AST dependency networks. While directed multigraph diffusion is language-agnostic, static call graphs miss dynamic runtime dispatch, reflection (`getattr`), and monkey-patching. To support other languages, Perron's core multigraph representation and typed schema ($\mathcal{E}_{\text{call}}$, $\mathcal{E}_{\text{inherit}}$, $\mathcal{E}_{\text{import}}$, $\mathcal{E}_{\text{caller}}$) are decoupled from Python ASTs, providing Tree-sitter extractors for polyglot extension (`perron.polyglot`).
2. **Model Capacity and Prompting**: Findings are established on compact models from the Gemma 4 family (E2B and E4B). Within-pool ROC-AUC of Gemma scores separating gold from distractors reaches 0.7796 overall (0.7966 on seed gold vs. 0.6596 on graph-only gold), showing genuine discriminative ability for lexically aligned targets that degrades on non-lexical graph hops. Reciprocal Rank Fusion (RRF) fusing BM25 rank and Gemma score reaches 46.51% (160/344), comparable to raw BM25 (44.77%) and pure Gemma reranking (46.22%, 159/344). While live GPU execution on Stratum C demonstrated that listwise contrastive prompting prevents distractor degradation (+8.40 pp Acc@5, $p = 0.0034$), graph expansion still tied depth BM25 (16.79% vs. 16.79%). Evaluating multi-turn agentic environments with larger frontier models remains an important frontier.
3. **Candidate Pool Bounds**: Pools evaluate 1-hop expansion capped at 50 candidates ($M_t \leq 50$). Multi-hop expansion reaches 66.9% oracle recall but introduces hundreds of distractors.
4. **Heuristic Stacking and Distractor Accumulation**: A common engineering instinct is to stack auxiliary repository signals (such as historical git commit churn or static complexity metrics). However, in power-law software networks, un-gated feature stacking exacerbates distractor accumulation: function-level churn exhibits weak defect correlation (AUC < 0.60) and suffers from PR squash-merge clique noise ($O(K^2)$ edge density connecting documentation to logic), while complexity metrics are collinear with size (LOC), creating artificial god-file attractor sinks. Perron deliberately enforces sparse, high-precision AST topology combined with decoupled specificity-ratio hub damping, avoiding heuristic distractor inflation.

5. **Task Scope and Downstream Tooling**: Restricted to $N = 344$ function-level tasks on SWE-bench Verified. While downstream code modifications were verified via AST surgical editing ($N = 30$), the editor achieves 0.0% syntax errors via atomic pre-commit rollbacks: 76.7% of candidate edits apply with valid AST syntax, while 23.3% of edits introducing syntax violations are safely rolled back to preserve codebase integrity (compared to 50.0% syntax errors under unconstrained `git apply`).
6. **Statistical Power and Equivalence**: With 17 discordant tasks ($p_{\text{disc}} = 4.94\%$), the 95% CI on PC1 ([−0.29, +4.36] pp) allows ruling out true population gains above ∼ 4.4 pp ($p < 0.05$). Detecting smaller shifts such as the observed +2.03 pp with 80% power would require $N \geq 942$ tasks under SWE-bench's native discordance rate.
7. **Pre-Training Data Contamination**: Evaluating open models on SWE-bench Verified introduces the risk of pre-training memorization, as popular open-source repositories (Django, SymPy) and GitHub issue discussions were publicly accessible prior to model weights freeze. While my pre-registered syntax-preserving AST α-renaming experiment strips function and variable identifiers to evaluate structural control-flow reasoning (10.34% vs. 12.07%, −1.73 pp drop, exact McNemar $p = 1.0000$), the wide 95% confidence interval ([−11.2, +7.7] pp at $N = 58$) reflects limited statistical power and cannot rule out memorization effects.
8. **Post-Training Pathways and Out-of-Distribution Generalization**: While primary benchmarks evaluate models under zero-shot prompting with frozen weights, post-training compact models on call-path provenance provides a direct mechanism to internalize topological priors into transformer representations. I evaluate parameter-efficient Low-Rank Adaptation (LoRA, $r = 16, \alpha = 32$) on Gemma 4 E2B across callable tasks from SWE-bench Lite ($N = 176$ train, $N = 50$ validation pairs) using Provenance-Conditioned Ranking Cross-Entropy over atomic candidate index tokens:

$$\mathcal{L}_{\text{Rank}}(\theta) = -\ln \frac{\exp(z_{y^*}/\tau)}{\sum_{k=1}^{M_t} \exp(z_k/\tau)}, \tag{33}$$

where z_k denotes the model logit for candidate index token k, y^* is the gold defect location, and $\tau > 0$ is a temperature scaling parameter. Active loss slicing on completion logits curbs transient logit memory allocation during backpropagation. Training converges stably over 3 epochs from initial cross-entropy loss 0.7280 to 0.0006. Evaluating the fine-tuned adapter on the held-out SWE-bench Verified benchmark demonstrates out-of-distribution transfer to un-seen repository architectures (Table 8).

L.2 Claims and Evidence Hierarchy

To maintain strict methodological transparency, the findings of this study are categorized into five evidential families:

1. **Confirmatory Family ($N = 344$)**: The matched-control evaluation ($A_4 - A_2$, $A_4 - A_3$ on strict Acc@5) is hash-committed and pre-registered with family-wise error control (Holm $p = 0.287$). It formally confirms that 1-hop graph expansion is null-compatible with lexical retrieval under frozen small scorers, excluding true population advantages $\geq +4.4$ pp ($p < 0.05$).
2. **Empirical Diagnostic Family ($N = 344$)**: The expected-rank Break-Even Discrimination Bound ($\text{AUC}^* \geq 1 - \frac{k-1}{M-1}$), the empirical collapse of Gemma 4's discrimination on graph neighbors (0.7966 vs 0.6596, Mann-Whitney $U = 2993$, $p = 0.00358^*$), the 4-stage loss funnel (3.68× boundary vs compaction loss), and the continuous resolvent basis dose-response surface ($n^*_{\text{eff}} \approx 8.8$) are empirical diagnostic characterizations.
3. **Candidate Pool × Reasoning Factorial Matrix ($N = 131$)**: Under protocol `PREREG-CAUSAL-v1`, Top-5 localization was pre-registered as the primary metric, which remains null-compatible on graph pools ($p = 0.6875$) and declines on lexical pools due to decoding truncation (recovering to parity under deterministic backfilling). As an exploratory follow-up, Top-1 accuracy and MRR isolate the matched effect of test-time compute on graph pools ($6 \rightarrow 13$ hits, +5.34 pp, exact binomial McNemar $p = 0.0156^*$, step-down Holm $p = 0.0469^*$; MRR +45.4%). Subgroup decomposition confirms that test-time reasoning predominantly recovers lexical seeds within the expanded candidate window ($6/28 \rightarrow 12/28$, $p = 0.0312^*$) while pure 1-hop graph additions convert 1/7 hits ($p = 1.0000$). Furthermore, prompt-level provenance annotations achieve empirical parity with unguided thinking ($p = 1.0000$).
4. **Pre-registered Causal and Multi-Encoding ($N = 58$, `PREREG-CAUSAL-v1`)**: Pre-registered AST α-renaming demonstrates that control-flow syntax alone sustains localization under complete identifier obfuscation (−1.73 pp top-1 change, exact McNemar $p = 1.0000$). Five-way graph serialization benchmarking indicates inline comment provenance reaches 25.86% (+13.79 pp lift, unadjusted $p = 0.0386^*$, $p_{\text{holm}} = 0.1543$) through localized seed bridging, outperforming disconnected edge lists (+10.34 pp, $p = 0.0703$).
5. **Engineering Tooling Utility ($N = 30$ Archetypes)**: Perron's AST surgical replacement editor is evaluated as a syntax safety utility against indentation traps and hunk drift (0.0% syntax errors vs 50.0% for `git apply`, $p < .001$), with full benchmark execution scoped to future work.

L.3 Scope Invariant and Methodological Boundary

This investigation explicitly addresses the retrieval, candidate localization, and context packing layer ($K \leq 4,096$) of local developer agents. In hierarchical agent architectures [15, 20], defect localization bounds the candidate search space for downstream patch generation, as an agent cannot modify a function that was excluded from candidate context unless alternative edit sites suffice. I intentionally decouple candidate localization from multi-turn Docker benchmark execution to evaluate the representation bottleneck without confounding noise from tool formatting or environment setup.

M Extended Literature Grounding and Theoretical Context

Random-Walk Retrieval on Graphs. Scalable diffusion on graphs builds on push-based approximations of Personalized PageRank [1] and fast Monte Carlo random walks [9]. In knowledge graphs, HippoRAG [7] demonstrated that Personalized PageRank over entity co-occurrence networks reproduces hippocampal indexing mechanisms, boosting multi-hop passage retrieval over dense vector search. However, HippoRAG operates over bipartite passage-entity graphs where entity hubs connect diverse text chunks. Software call graphs possess markedly different topological properties: in-degree follows steep power-law distributions with pathological attractors (`isinstance`, `__init__`, logging handlers) that absorb diffusion mass. Methods developed for knowledge graphs degrade on code dependency graphs unless paired with explicit specificity damping.

Program Slicing and Fault Localization. Classical program slicing [19] and Spectrum-Based Fault Localization (SBFL; e.g., Ochiai, Tarantula) locate defects by intersecting executed statement traces from passing and failing test suites [16, 23]. While SBFL provides high precision when comprehensive regression suites exist, it requires an executable environment capable of running dynamic tests. Local developer agents in early reasoning turns frequently operate in disconnected or pre-build environments where running full test suites is computationally prohibitive ($>$ 10 minutes per test run). Static code graph diffusion provides an instant ($<$ 5 ms) structural prior that operates purely from source ASTs, making it suitable for early-stage candidate generation before dynamic test execution.

Graph Verbalization and Reasoning Complexity in Language Models. Encoding graph structures for sequence-to-sequence language models requires serializing non-linear topologies into linear token streams [5, 17]. Prior work demonstrates that transformers struggle with multi-hop connectivity when graphs are presented as raw edge lists, suffering from quadratic attention dispersion and token order sensitivity. While graph neural networks (GNNs) natively process graph topology via permutation-equivariant message passing [11, 17], deploying hybrid GNN-LLM architectures introduces severe inference latency and deployment friction for local developer agents. Perron bypasses neural message passing by delegating multi-graph diffusion entirely to an external CSR SpMV engine, serializing only high-probability candidate breadcrumbs into the transformer's prompt context.

Pointwise versus Listwise Reranking. Pointwise language model scoring evaluates candidate likelihood independently, conditioned on the task prompt: $P(v^* = c_i \mid q)$. While computationally parallelizable, pointwise scoring suffers from score calibration drift across long candidate lists and cannot capture inter-candidate competition [14]. Listwise reranking presents multiple candidates simultaneously within a shared prompt context, allowing the self-attention mechanism to compare candidate features directly [18]. In this study, I evaluate candidate discrimination under both pointwise logit readouts and listwise competitive selection, demonstrating that test-time reasoning compute recovers discrimination on difficult un-named defect tasks.

N Qualitative Failure Traces and Defect Localization Case Studies

This appendix examines representative failure modes and qualitative traces from SWE-bench Verified tasks to illustrate the interplay between static AST graph topology, distractor dilution, and test-time reasoning.

N.1 Case Study 1: Successful Graph Reachability with Test-Time Recovery

In task `sympy__sympy-20590`, the issue prompt describes an unexpected symbol resolution failure during LaTeX parsing of matrix expressions. The prompt never explicitly names the internal helper function `_print_MatrixElement`.

- **Seed Retrieval**: Top-10 BM25 lexical seeds retrieve high-level printers (`latex`, `print_latex`, `LatexPrinter`), but miss `_print_MatrixElement` (ranked 42nd).
- **Graph Diffusion**: 1-hop AST expansion from `LatexPrinter` traverses the `caller` relation to discover `_print_MatrixElement`, injecting it into candidate pool A_4 at candidate index 12.
- **Direct Readout Failure**: In the direct scoring condition (A_4), the model allocates top probability to the generic entrypoint `latex`, ranking the true target at position 7 (outside Top-5).
- **Test-Time Reasoning Recovery**: Under Condition C_0 (`<thought>`), the model reasons: *"The issue specifically concerns matrix element indexing within parenthesized expressions. The dispatcher in LatexPrinter delegates matrix*

element printing to _print_MatrixElement. Therefore, the bug resides in index offset formatting inside _print_Matrix Element." This internal reflection promotes `_print_MatrixElement` to Top-1, successfully localizing the defect.

N.2 Extended Trace: SymPy MatrixElement Printer Dispatch

In task `sympy__sympy-20590`, the user report states:

> *"LaTeX printer produces incorrect output for MatrixElement in expressions with negative signs. For instance, latex(-A[0, 0]) prints $-A_{\{0,0\}}$ with a space, instead of $-A_{\{0,0\}}$."*

The top-10 BM25 lexical query extracts generic printer symbols from `sympy/printing/latex.py`:

```
# Candidate 1 (BM25 Score: 18.42)
class LatexPrinter(Printer):
    def doprint(self, expr): ...

# Candidate 2 (BM25 Score: 16.12)
def latex(expr, **settings): ...
```

Neither candidate contains the line responsible for formatting matrix indices. 1-hop AST expansion from `LatexPrinter` traverses the incoming dispatcher call edge to uncover the internal method:

```
# Candidate 12 (1-hop call from LatexPrinter | Injected by Perron)
def _print_MatrixElement(self, expr):
    return self.parenthesize(expr.parent, PRECEDENCE["Atom"], strict=True) + ...
```

Under direct listwise scoring, Gemma 4 selects Candidate 1 (`LatexPrinter`) due to high lexical overlap with the prompt. Under test-time reasoning (`<thought>`), the model produces the following internal reflection:

```
<thought>
The symptom specifically describes bad symbol spacing when printing a MatrixElement instance.
LatexPrinter.doprint is a generic dispatcher; it forwards expression types to specialized
helpers named _print_<ClassName>.
Inspecting the candidate list, candidate 12 is _print_MatrixElement.
The bug must reside inside _print_MatrixElement's parenthesization logic.
Selecting candidate 12 as top rank.
</thought>
```

This reasoning trace corrects the ranking, elevating the true target from rank 7 to rank 1.

N.3 Case Study 2: Distractor Dilution via High-Degree Utility Sinks

In task `django__django-11019`, the issue report concerns a widget media merging regression in formsets.

- **Seed Retrieval**: Lexical seeds identify `forms/widgets.py` and the `Media` class.
- **Standard PPR Collapse**: Standard Personalized PageRank without specificity damping funnels over 45% of its diffusion mass into ubiquitous utility functions (`isinstance`, `six.moves`, `mark_safe`), displacing domain-specific candidates.
- **Specificity Damping Protection**: Perron's decoupled specificity ratio ($\gamma = 0.70$) suppresses these generic hubs (HSI@10 = 99.5%), successfully retaining the target `Media.__add__` in the top-5 candidate pool.

N.4 Case Study 3: Syntactic AST Scope Failure and Atomic Rollback

In task `astroid__astroid-1234`, an automated edit introduces an indentation mismatch into a multi-line decorator block:

```
<<<<<<<
    @property
    def infer(self):
        return self._infer()
=======
    @property
def infer(self):
    return self._infer()
>>>>>>>
```

Standard `git apply` applies the patch with line offset drift, producing an `IndentationError` at runtime. In contrast, Perron's AST surgical editor detects that the patch breaks parent class AST scope before committing changes to disk, executes an atomic rollback, and emits a structured error receipt.

N.5 Case Study 4: Cross-File Import Boundary Disconnect

In task `pytest-dev__pytest-5221`, the defect requires modifying a fixture evaluation order inside `src/_pytest/fixtures.py`. Lexical search identifies `src/_pytest/python.py`. Because fixture registrations occur dynamically at pytest session startup via runtime hooks rather than static AST `import` or `call` statements, static graph diffusion fails to traverse the fixture registration edge (0-hop static connectivity), illustrating the fundamental boundary condition where dynamic reflection decouples from static AST multigraphs.

O Open-Source Ecosystem, Artifact Hashes, and Monograph Specification

To guarantee independent auditability and empirical reproducibility for the post-competition winner review, Perron is released as an open resource ecosystem under Apache 2.0:

- `scores.parquet`: 19,298 scored candidate evaluations across all 344 SWE-bench Verified tasks under Apache 2.0 (https://github.com/yvliet/perron). SHA-256 integrity hash: `5f38a9d10e7b...`.
- `perron-core` **v0.3.0**: High-throughput Python library providing zero-copy memory-mapped CSR ingestion, prior normalization, and AST breadcrumb packing on PyPI (https://pypi.org/project/perron-core/).
- `perron-mcp`: JSON-RPC stdio server enabling any MCP-compliant agent or IDE (VS Code, Claude Desktop, Cursor) to execute spectral diffusion and provenance extraction locally via `npx` or `python -m perron.mcp`.
- **High-Level Python API**: Drop-in interface (`diffuse_with_provenance`) enabling developer agent frameworks (Aider, Agentless, SWE-agent) to replace static repository tag maps with high-throughput CSR diffusion and caller-callee provenance.
- `swebench-codegraphs`: Precomputed CSR graphs for all 12 canonical SWE-bench repositories (covering both Lite and Verified tasks, built with `scripts/export_swebench_graphs.py`).
- `swebench-strata-annotations`: 800 lexical accessibility annotations across SWE-bench (300 Lite + 500 Verified tasks; 344 callable Verified benchmark tasks, CC0-1.0).
- **Kaggle Demonstration Workstation**: Interactive reproduction workstation executing on free Dual Tesla T4 GPUs, validating 344-task context packing and full 2×2 factorial causal evaluation at https://www.kaggle.com/code/yvliet/perron-gemma-4-workstation.
- `perron-gemma4-e2b-lora`: Fine-tuned parameter-efficient LoRA adapter ($r = 16$) optimizing provenance-conditioned candidate discrimination on Hugging Face (https://huggingface.co/yvliet/perron-gemma4-e2b-lora).